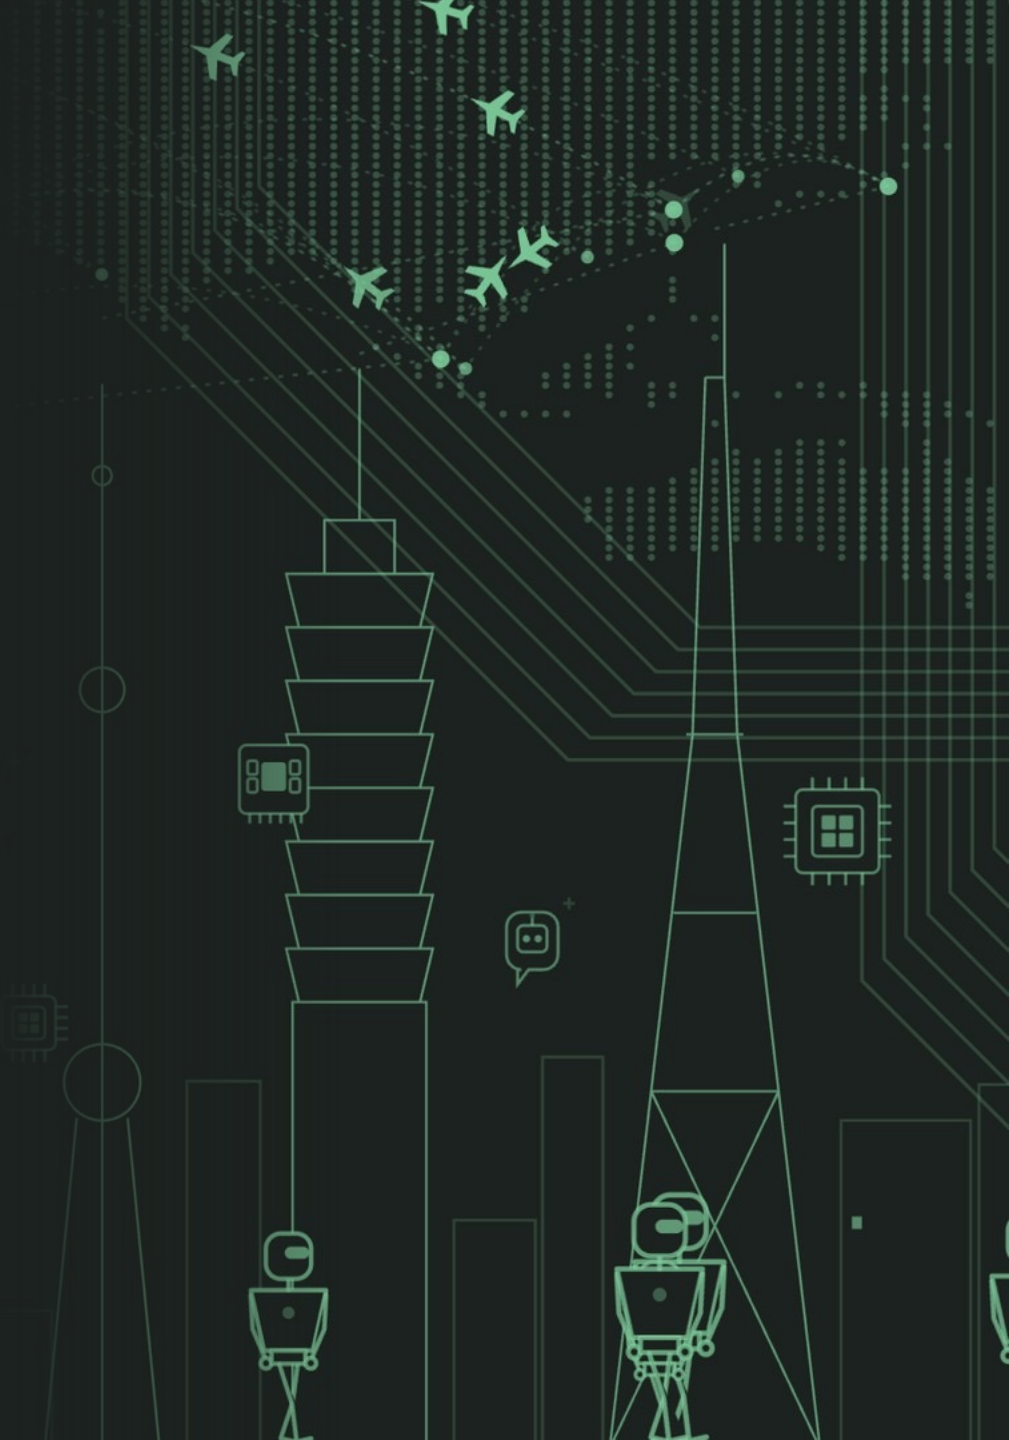
ARC Research · AI Update · Nov 12, 2026

NVIDIA after Jensen

The founder, the news, the interim CEO and the balance sheet

Founder-CEO era	Revenue, Q2 FY27	Cash and securities	Stock since Oct 29
33 years	$96.2B	$99.4B	−23%



Data as of Nov 12, 2026 · Educational purposes only, not investment advice

Agenda

01	**Current news**	What happened, who is investigating, the market	Slides 3–8
02	**Company background**	What NVIDIA is and what it sells	Slides 9–11
03	**The founder**	Jensen Huang and the company he built	Slides 12–20
04	**Financials**	Margins and the balance sheet	Slides 21–27
05	**The interim CEO**	Colette Kress and her background	Slides 28–36
06	**The Apple precedent**	Apple in 2011: Jobs, Cook, demand and rivals	Slides 37–47
07	**The competition**	Google, Microsoft, Amazon and AMD chips	Slides 48–52
08	**Shoes to fill**	TSMC and Korea’s HBM makers	Slides 53–56
09	**ARC Research view**	Our view on what comes next	Slides 57–67

Also: Sources (slides 68–69)

Section 1 of 9

Current news

What is happening, in plain terms

COMING UP

4 What is happening, in plain terms

5 How it unfolded, day by day

6 Who is investigating what

7 The share price

8 Market reaction and what to watch



What is happening, in plain terms

THE CEO

Jensen Huang

Admits an affair, Nov 8; denies any security breach

Reported relationship →

THE EXECUTIVE

A Megaspeed executive

Described in reports as a suspected Chinese intelligence asset

Works at →

THE CUSTOMER

Megaspeed, Singapore

Bought $4.8B+ of NVIDIA chips since mid-2024

Alleged →

THE ISSUE

Chips to China?

DOJ criminal probe of export breaches

THEN

A second relationship, with a direct report, is reported Nov 11

RESULT

Board accepts his resignation; announced Nov 12

WHY IT MATTERS

Export-control risk for NVIDIA, not just a scandal

Allegations and press reports, not findings. Megaspeed says it will cooperate.

Fourteen days to a resignation

Oct 29
First report ties Huang to an alleged Chinese intelligence asset

Nov 5
Three House committee chairs ask for a briefing

Nov 8
Huang admits infidelity; board opens an investigation

Nov 10
Board retains Joon Kim and Wachtell Lipton

Nov 11
Second relationship reported; board accepts resignation

Nov 12
Exit announced before the open

THE NEW LINE-UP



Jensen Huang
Out as President, CEO and director



Colette Kress
Interim CEO; stays EVP and CFO



Harvey C. Jones
First-ever board chair; ex-Synopsys CEO

Who is investigating what

NVIDIA BOARD

Independent review

Joon Kim and Wachtell Lipton, hired Nov 10

JUSTICE DEPT.

Criminal probe

Megaspeed’s alleged chip smuggling to China

CONGRESS

Huang to testify

Before Thanksgiving, reportedly closed-door

SENATE

Records demanded

Warren and Coons: export and contact records

COMMERCE

Lutnick under fire

Bipartisan calls to explain lapsed controls

INVESTOR RISK

Export breach

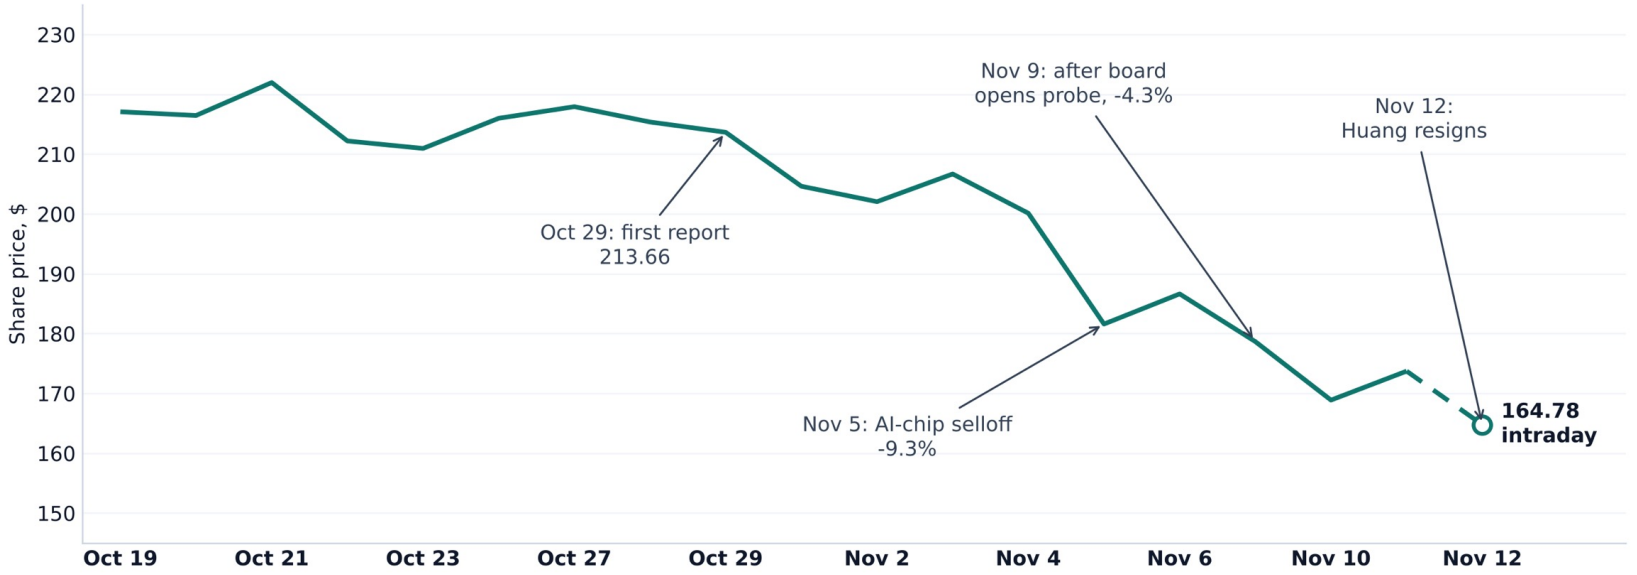
Penalties or tighter China rules if chips leaked

The stock: down 23% since the first report

ARC RESEARCH | AI UPDATE | NVIDIA SHARE PRICE

NVIDIA shares: down 23% since the first report

Daily closes, Oct 19 to Nov 12, 2026. Record close was $239.24 on Oct 6; Nov 12 is about 31% below it



Source: Investing.com historical prices (daily closes, $). Nov 12 point = last price in that table, $164.78, taken during the session, not an official close.
Record close per StockTitan and Yahoo Finance. -23% = 164.78 / 213.66 - 1 (from the Oct 29 close); -31% = 164.78 / 239.24 - 1; ARC Research calculations.
Not all of the fall is the scandal: Nov 5 was an AI-chip selloff (Trefis).

SO WHAT

$213.66 to about $165 in two weeks

31% below the Oct 6 record

Nov 5 drop was an AI-chip selloff too

The market’s verdict, and what’s next

−23%
Since the Oct 29 report

−31%
From the Oct 6 record close

−$1.4T
Market value in two weeks, per Axios

<$4T
Market cap, first time since 2025

WHAT TO WATCH

Nov 17
Q3 FY27 earnings; guide $108B ±2%

CEO search
Internal and external candidates

Congress
Huang agrees to testify, per CNBC

Anthropic IPO
NVIDIA pledge of up to $10B

Section 2 of 9

Company background

What NVIDIA is, and where its money comes from

COMING UP

10 NVIDIA at a glance

11 Revenue by platform



NVIDIA at a glance

1993
Founded; based in Santa Clara

$215.9B
Fiscal 2026 revenue, +65%

75.0%
Gross margin, Q2 FY27

92%
Of revenue from Data Center

WHAT IT SELLS

AI systems
Blackwell Ultra now; Vera Rubin in full production

Networking
Links GPUs into AI factories

CUDA software
The developer platform since 2006

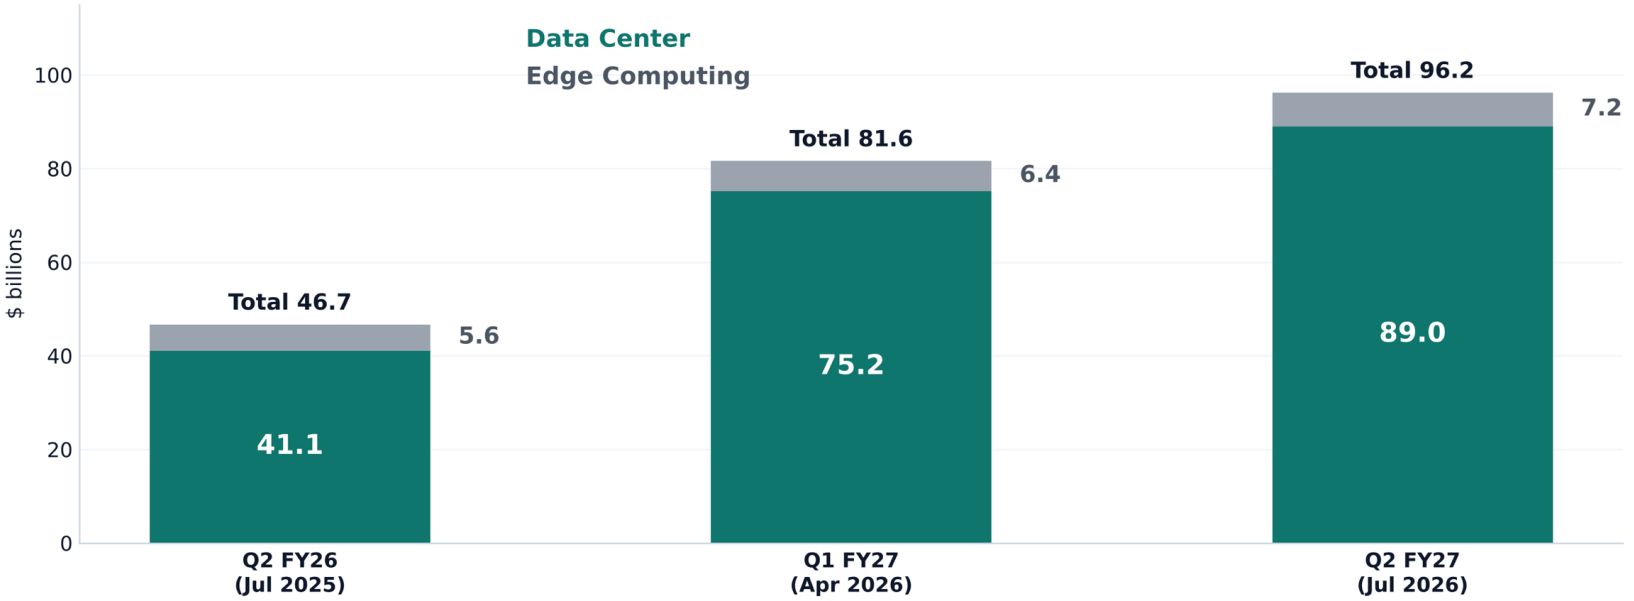
Fabless
Designs the chips; TSMC builds them

Where the money comes from: Data Center

ARC RESEARCH | AI UPDATE | NVIDIA REVENUE BY PLATFORM

Data Center is 92% of NVIDIA's revenue

Quarterly revenue by market platform: Data Center $89.0B of $96.2B in Q2 FY27



Source: NVIDIA Q2 FY27 Form 10-Q and results releases (Form 8-K). $ billions. Data Center as reported (recast platform view).
Edge Computing: Q2 FY27 $7.2B as reported; earlier quarters derived as total revenue less Data Center (46.7 - 41.1; 81.6 - 75.2).
92% = 89.0 / 96.2, ARC Research calculation.

SO WHAT

Data Center $89.0B, +117% y/y

Edge Computing $7.2B, +27% y/y

No China data-center sales in guidance

Section 3 of 9

The founder

Who Jensen Huang is, and what 33 years built

COMING UP

13	Profile and milestones
14	Four skills: technical, financial, political, deals
15	The dealmaker: his deal book
16	Mellanox: the doubted deal that paid off
17	AI-pilled for over a decade: the bets
18	The open-source bet
19	How the CUDA bet played out
20	The growth, in one chart



Jensen Huang: 33 years at the helm



Jensen Huang

Co-founder · CEO 1993–2026

1993	Founded at a Denny’s in San Jose
1999	IPO; GeForce 256, marketed as the first GPU
2006	CUDA opens the GPU to general computing
2012	AlexNet trains on NVIDIA GPUs
2022–26	Generative AI boom; a $4 trillion company
Nov 2026	**Resigns as President, CEO and director**

BORN

1963, Tainan, Taiwan

EDUCATION

BSEE Oregon State · MSEE Stanford

BEFORE NVIDIA

LSI Logic, AMD

CO-FOUNDERS

Chris Malachowsky, Curtis Priem

Four skills in one person

TECHNICAL

The engineer

- Chip designer at LSI Logic and AMD
- AI-pilled since 2012: GPU, CUDA, DGX bets
- A new platform every year; Vera Rubin in production

FINANCIAL

The capital allocator

- Balance sheet as a financing engine
- $30B OpenAI stake; supplier capacity pre-booked
- 75% gross margin at $96B a quarter

POLITICAL

The diplomat

- Lobbied Trump to reopen China chip sales in 2025
- Sold governments on sovereign AI
- Dealt directly with Korea’s chairmen

DEALS

The dealmaker

- $30B OpenAI stake; up to $10B in Anthropic
- $500B chip-financing pact with Wall Street
- Mellanox bought; Arm attempted

THE IRONY

Politics and China are where his tenure ended

THE GAP

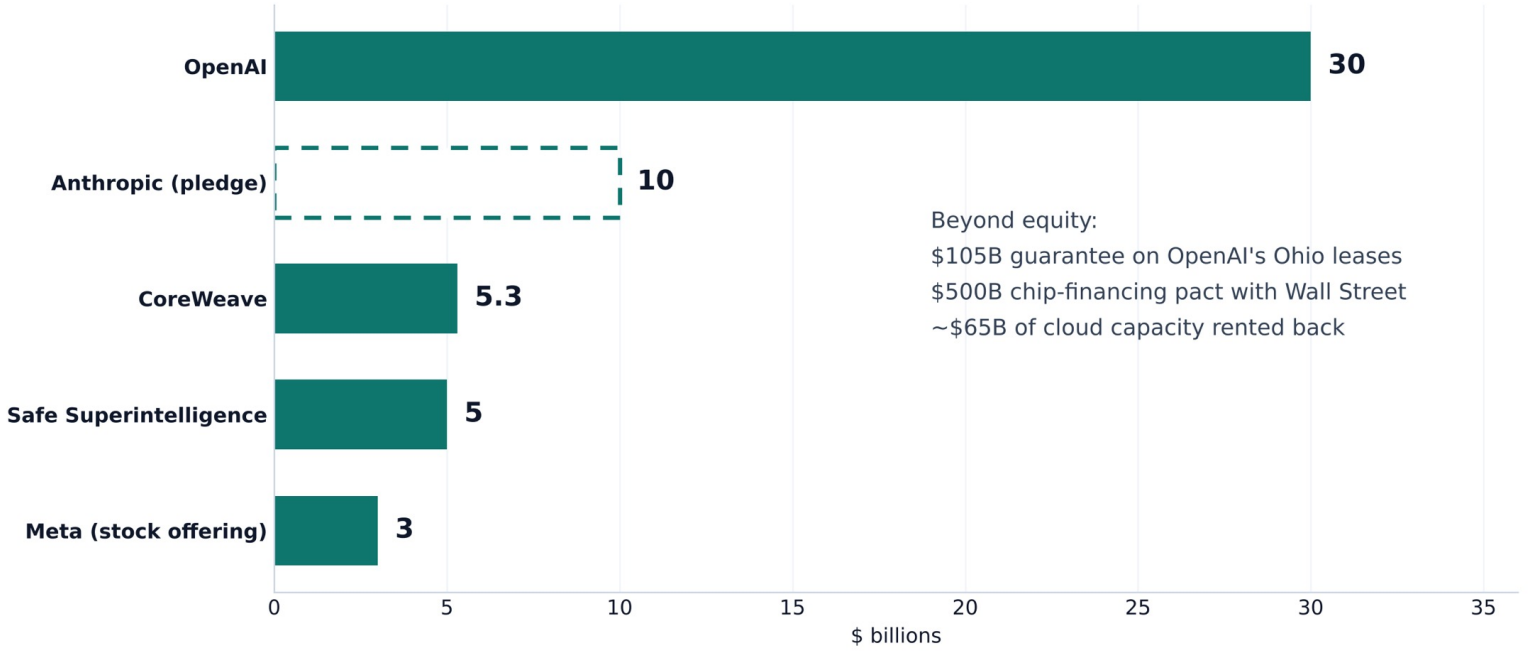
A successor must cover all four; Kress is strongest on finance

The dealmaker: Huang’s deal book

ARC RESEARCH | AI UPDATE | JENSEN THE DEALMAKER

Huang's deal book: equity in the companies that buy his chips

Largest equity commitments in AI labs and clouds, $ billions, plus the guarantees and financing behind them



Source: Yahoo Finance (Nov 12, 2026) tally of NVIDIA's deals and 10-Q; CNBC (Aug 18, 2026) on OpenAI and the Wall Street financing pact; NVIDIA Q2 FY27 call via press. $ billions. Solid = invested; dashed = pledged, not yet funded (Anthropic, ahead of its IPO). CoreWeave = $3.3B stake + $2B more. Press-reported figures.

SO WHAT

Stakes in the buyers of his chips

Guarantees and financing on top

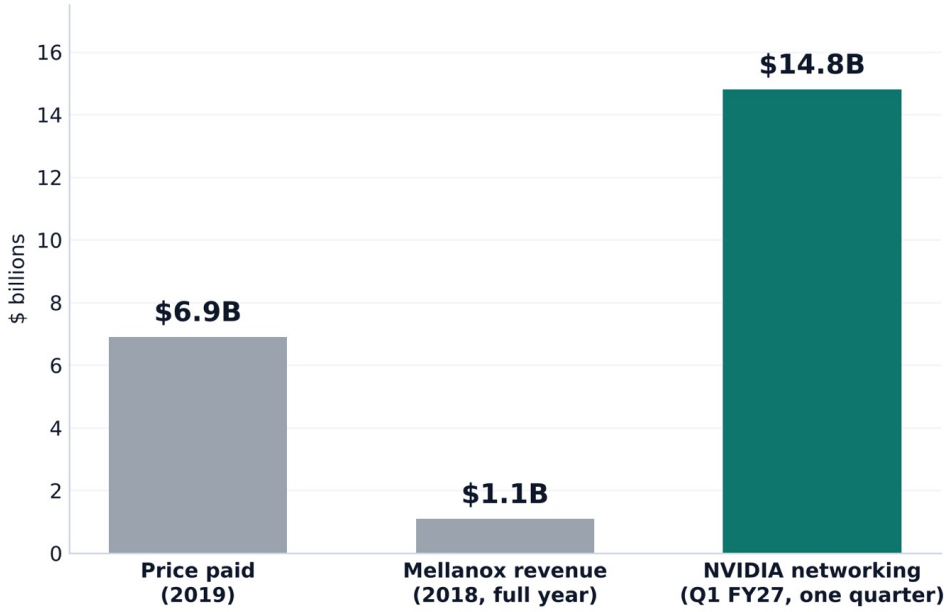
Mellanox bought; Arm blocked in 2022

The Mellanox bet: doubted in 2019, now core

ARC RESEARCH | AI UPDATE | THE MELLANOX BET

Doubted in 2019, Mellanox now sells twice its price every quarter

NVIDIA paid $6.9B for a $1.1B-revenue networking firm; that business now sells $14.8B in a single quarter



WHAT CRITICS SAID IN 2019

A "head scratcher"

Bloomberg Intelligence: little overlap with NVIDIA

Not a game changer

RBC: only ~$1.1B of revenue to add

Paid up

Morningstar: ~56% over pre-rumour price

Huang's view

Turns NVIDIA into a data-center-scale company

Source: deal announced Mar 11, 2019 at $125 a share, $6.9B in cash, a 14% premium (Bloomberg); closed April 2020. Mellanox 2018 revenue $1.09B (Morningstar; Bloomberg). NVIDIA Data Center networking revenue $14.8B in Q1 FY27 (quarter to Apr 26, 2026), up 199%, NVIDIA results release; in Q2 FY27 networking grew a further 18% q/q (earnings call). Critics' views as reported at the time: Bloomberg Intelligence via Benzinga; RBC Capital Markets via Benzinga; Morningstar. 14.8 / 6.9 = 2.1x, ARC Research calculation.

SO WHAT

Paid $6.9B in 2019; critics doubted it

Networking: $14.8B in one quarter

Made NVIDIA a systems company

AI-pilled for over a decade: Jensen’s bets

1999	**The GPU.** GeForce 256: graphics gets its own processor
2006	**CUDA.** Every GPU made programmable, years before a market
2012	**Deep learning.** After AlexNet, the company turns to AI
2016	**AI systems.** First DGX-1 supercomputer, delivered to OpenAI
2019–20	**Networking.** Buys Mellanox for about $6.9B
2022–26	**The payoff.** Data Center revenue $15B to about $194B

THE PATTERN
Invest years before the market exists

THE COST
CUDA was carried for years with little return

THE RESULT
92% of revenue is now AI data centers

The open-source bet: models for everyone

Mar 2025	**GR00T N1.** Open robotics foundation model at GTC
Mar 2026	**Nemotron 3 Super** (120B), then the Nemotron Coalition at GTC
Apr–Jun	**Nano Omni** (30B, multimodal) and **Ultra** (~550B)
Aug 2026	**Nemotron 3.5 Lightning** for fast agents
Nov 10	**OpenShell donated;** wider open-source commitment

WHY

Closed labs build their own chips; open models run on NVIDIA

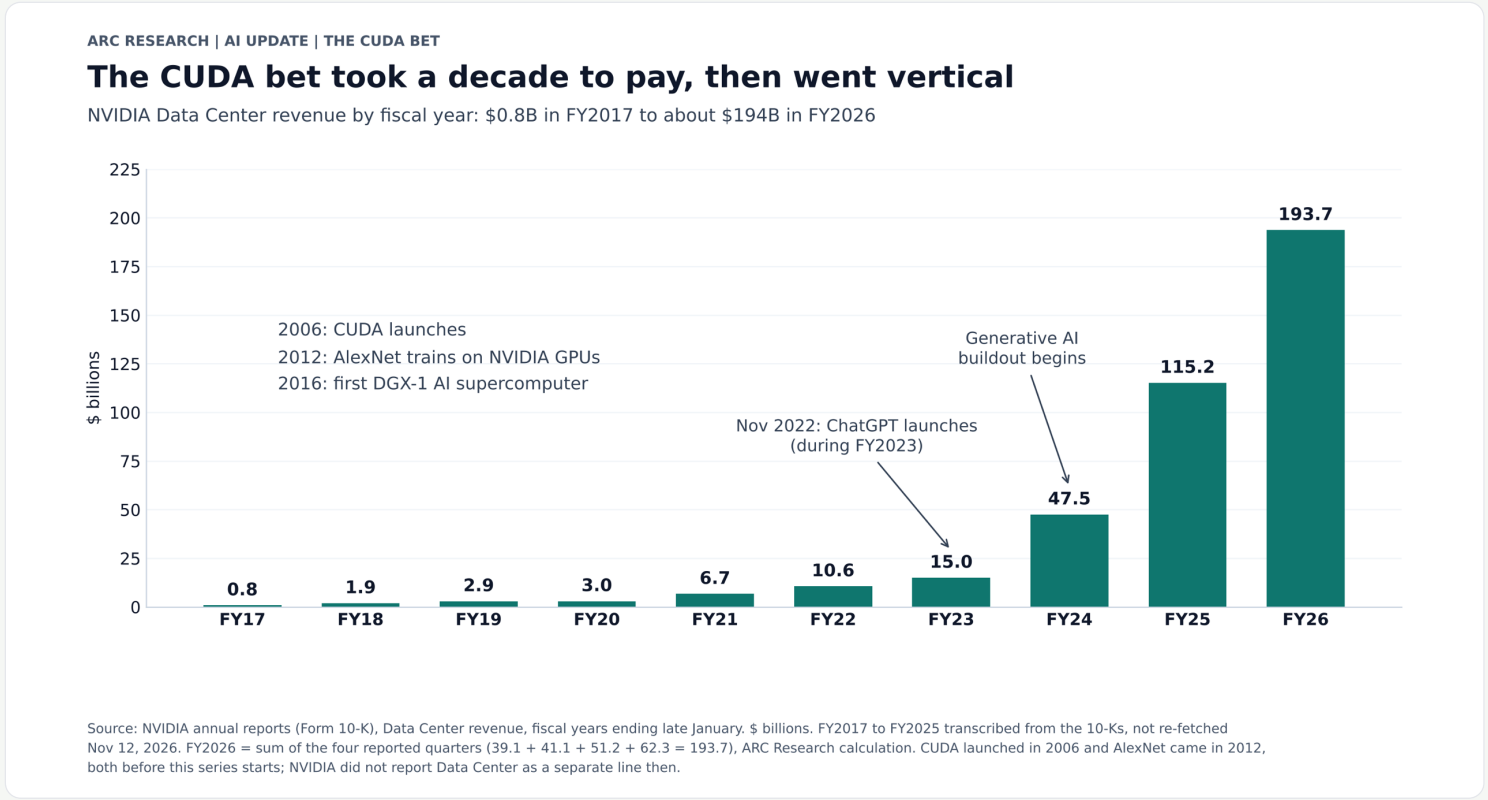
THE EFFECT

More customers, and a wider CUDA moat

THE LIMIT

The models are open; CUDA stays proprietary

The CUDA bet: ten years of patience, then lift-off



AUG 2016: THE FIRST DGX-1



Huang hand-delivers it to OpenAI

2006 CUDA · 2012 AlexNet · 2016 DGX-1

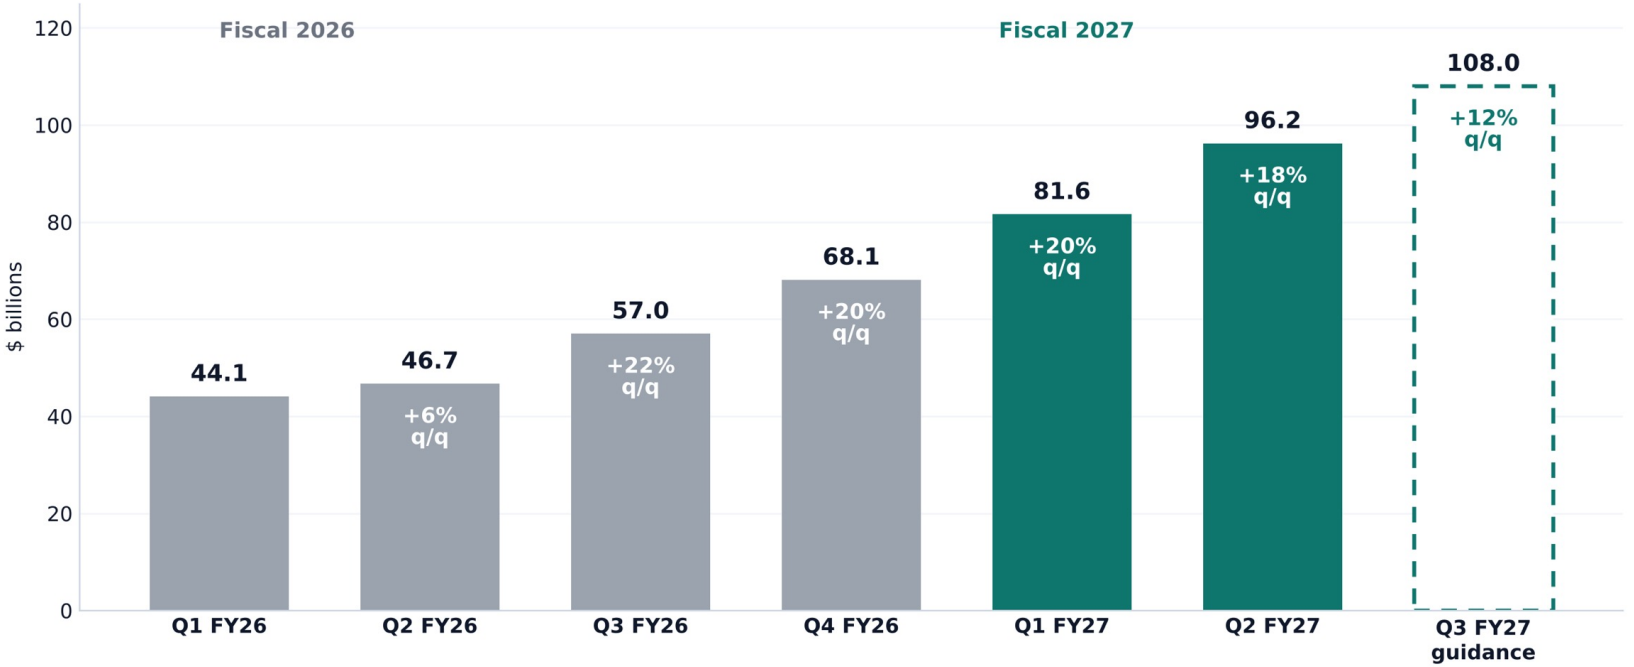
FY26: ~$194B Data Center revenue

What he built: revenue doubled in a year

ARC RESEARCH | AI UPDATE | NVIDIA QUARTERLY REVENUE

NVIDIA quarterly revenue more than doubled in a year

Quarter-on-quarter growth of 18-22% in each of the last four quarters; Q3 FY27 guidance $108.0B ±2% (dashed) implies about +12%



Source: NVIDIA quarterly results releases (Form 8-K exhibits), Q1 FY26 to Q2 FY27. $ billions, fiscal quarters (FY27 ends Jan 2027).
Q1 FY26 derived: six-month revenue $90,805M less Q2 $46,743M. Dashed bar = company guidance, not a reported result. Q/q growth from $M figures (44,062; 46,743; 57,006; 68,127; 81,615; 96,221), ARC Research calculation; guidance growth = 108,000 / 96,221 - 1 at the midpoint.

SO WHAT

$96.2B last quarter, +106% y/y

+18–22% q/q for four quarters

Guide $108B: about +12% q/q

Section 4 of 9

Financials

Margins, a fortress balance sheet, and how NVIDIA puts it to work

COMING UP

22 Revenue and margins by quarter

23 Revenue and P/E vs other chipmakers

24 How strong the balance sheet is

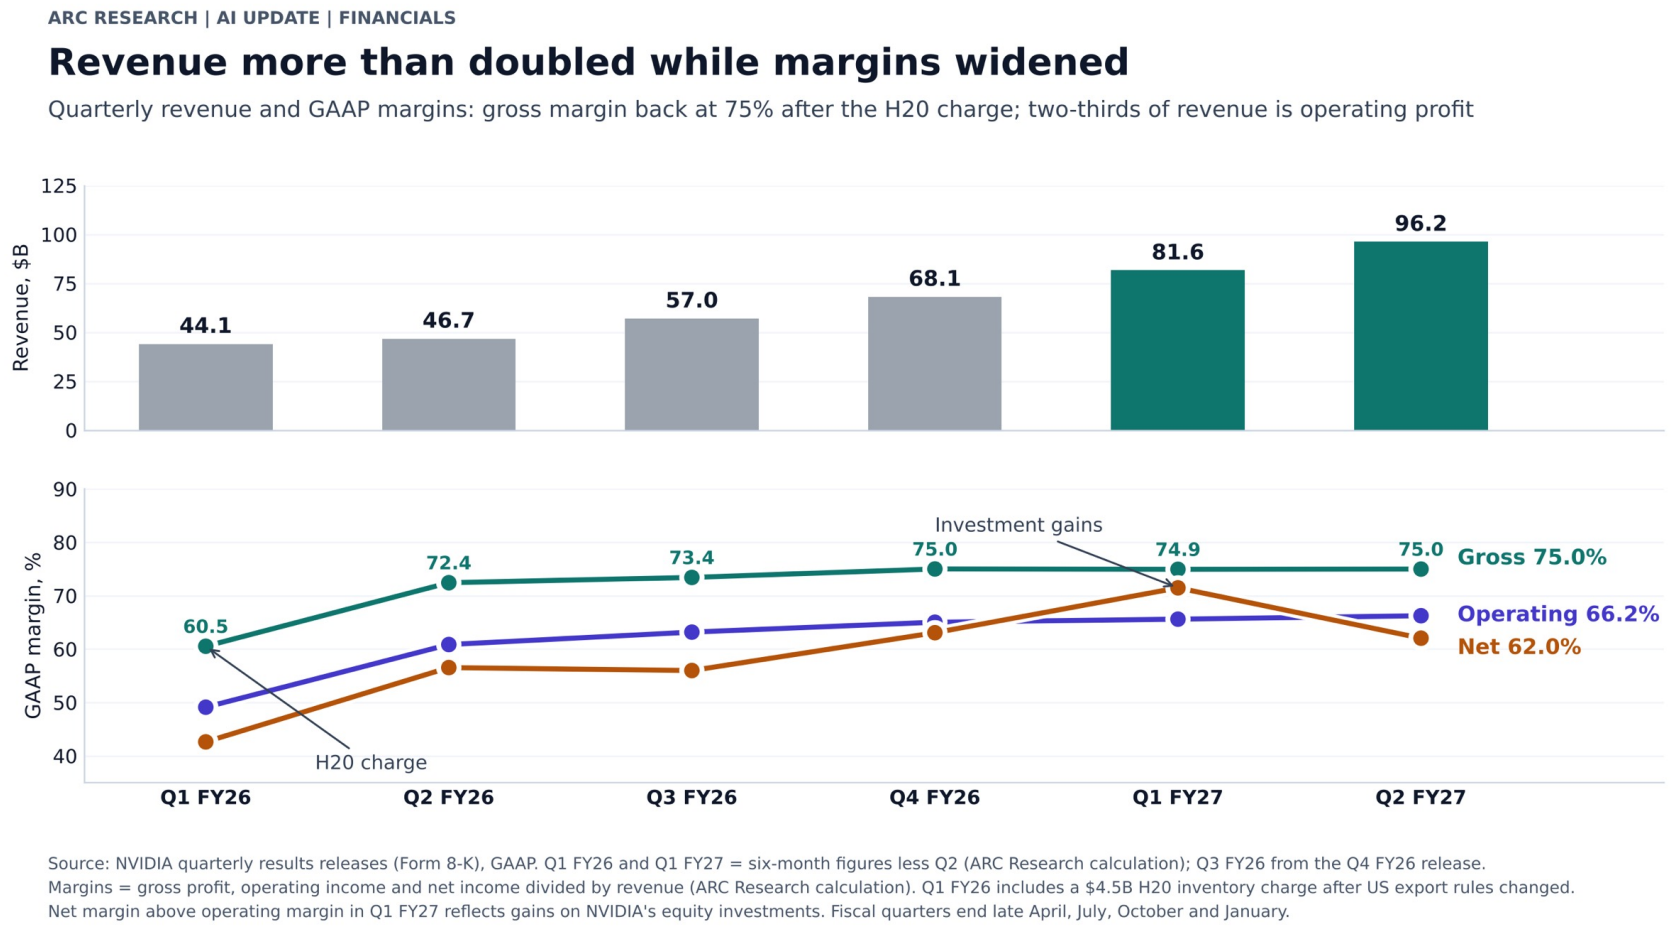
25 The financing flywheel

26 The venture arm: $94B of stakes

27 $530B of commitments



Revenue doubled; margins widened



Source: NVIDIA quarterly results releases (Form 8-K), GAAP. Q1 FY26 and Q1 FY27 = six-month figures less Q2 (ARC Research calculation); Q3 FY26 from the Q4 FY26 release.
Margins = gross profit, operating income and net income divided by revenue (ARC Research calculation). Q1 FY26 includes a $4.5B H20 inventory charge after US export rules changed.
Net margin above operating margin in Q1 FY27 reflects gains on NVIDIA's equity investments. Fiscal quarters end late April, July, October and January.

SO WHAT

Revenue $44B to $96B in five quarters

Gross margin 75%; operating 66%

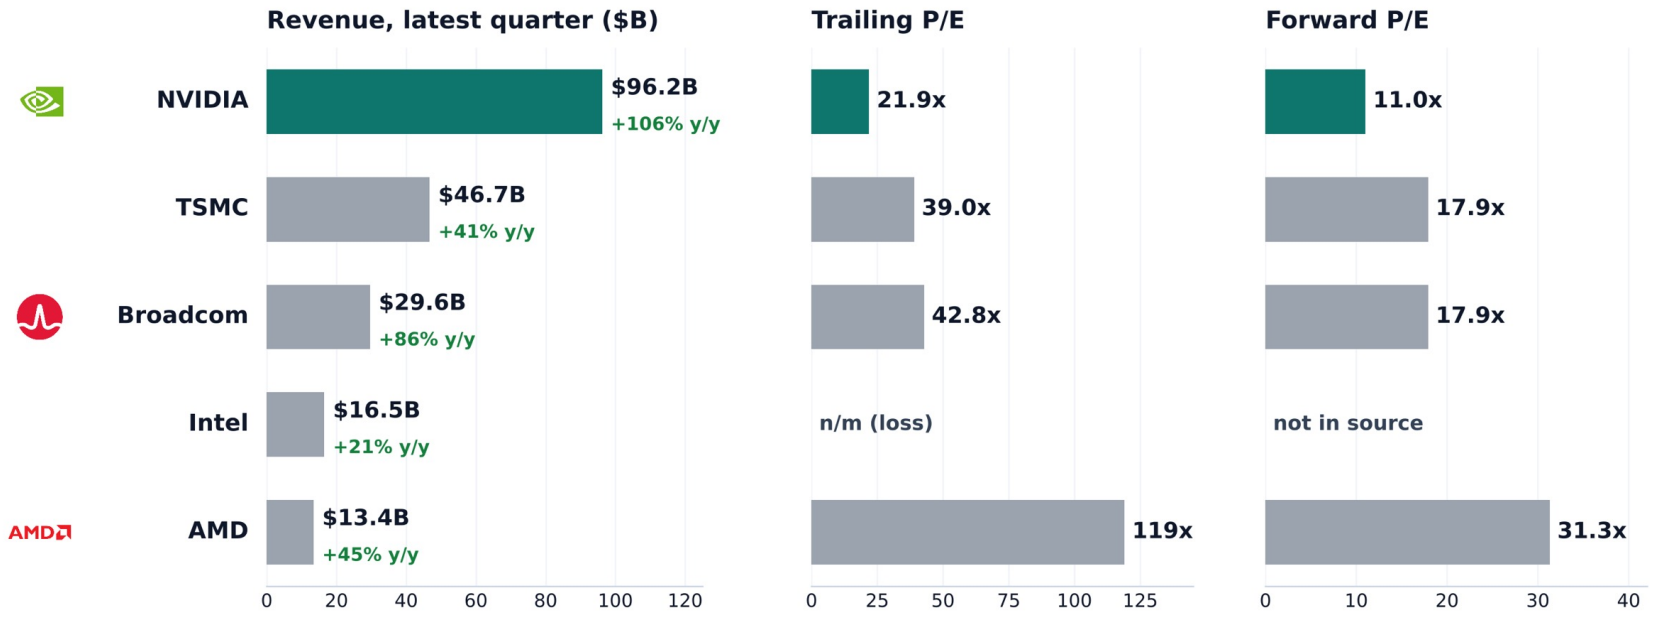
Q1 FY26 dip: $4.5B H20 charge

Most revenue, lowest P/E among chipmakers

ARC RESEARCH | AI UPDATE | FINANCIALS

NVIDIA has the most revenue, yet trades at the lowest P/E

Latest quarterly revenue and price-to-earnings ratios of the largest AI chipmakers; investors pay less per dollar of NVIDIA profit



Source: revenue and y/y growth from company results: NVIDIA Q2 FY27 (to Jul 26, 2026); TSMC Q3 2026 (US$, 6-K); Broadcom Q3 FY26 (to Aug 2); Intel Q3 2026; AMD Q3 2026. Quarters differ. Trailing P/E: CompaniesMarketCap, Nov 12, 2026 (price / last four quarters' EPS); Intel's is negative (trailing loss), shown as n/m. Forward P/E: Finviz via InvestingLive, Nov 12 (next-year consensus EPS; Intel not given). Third-party data, not company figures; ratios move with the share price daily.

SO WHAT

Biggest revenue, fastest big grower

Lowest P/E: 22x trailing, 11x forward

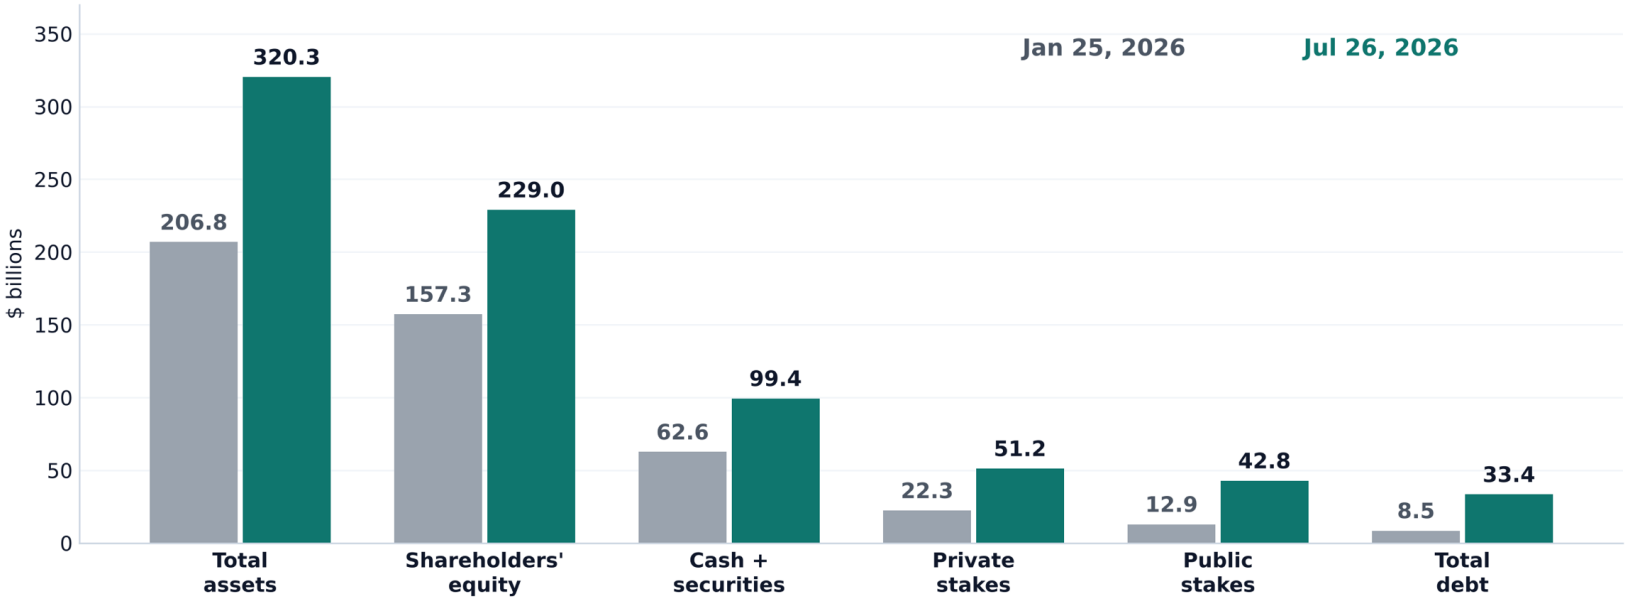
Market discounts the edge's durability

A fortress: $66B more cash than debt

ARC RESEARCH | AI UPDATE | NVIDIA BALANCE SHEET

A fortress: $99B of cash and securities against $33B of debt

Six months to Jul 26, 2026: total assets +55%, and both equity-stake lines more than doubled



Source: NVIDIA Q2 FY27 results release (Form 8-K, Aug 26, 2026), condensed balance sheets. $ billions.
Cash + securities = cash, marketable debt and marketable equity (public stakes, shown again separately). Private stakes = non-marketable securities.
Debt = short-term + long-term. +55% = 320.3 / 206.8 - 1, ARC Research calculation.

SO WHAT

$69.9B free cash flow, first half

$26.0B returned to holders in Q2

Debt rose to $33.4B from $8.5B

The balance sheet as a strategic tool

NVIDIA CAPITAL

$99.4B

Cash and securities; $69.9B first-half free cash flow

Invests, guarantees →

AI LABS AND CLOUDS

- OpenAI: $30B stake
- Anthropic: up to $10B
- Safe Superintelligence: $5B
- Meta: $3B of its offering
- CoreWeave: $3.3B, plus $2B

Buy chips →

NVIDIA REVENUE

$96.2B

Q2 FY27, +106% y/y

← **The loop closes:** chip revenue becomes cash flow that refills the balance sheet

PUBLIC STAKES

$42.8B

From $12.9B in January

PRIVATE STAKES

$51.2B

From $22.3B in January

“a scarcity of capital”

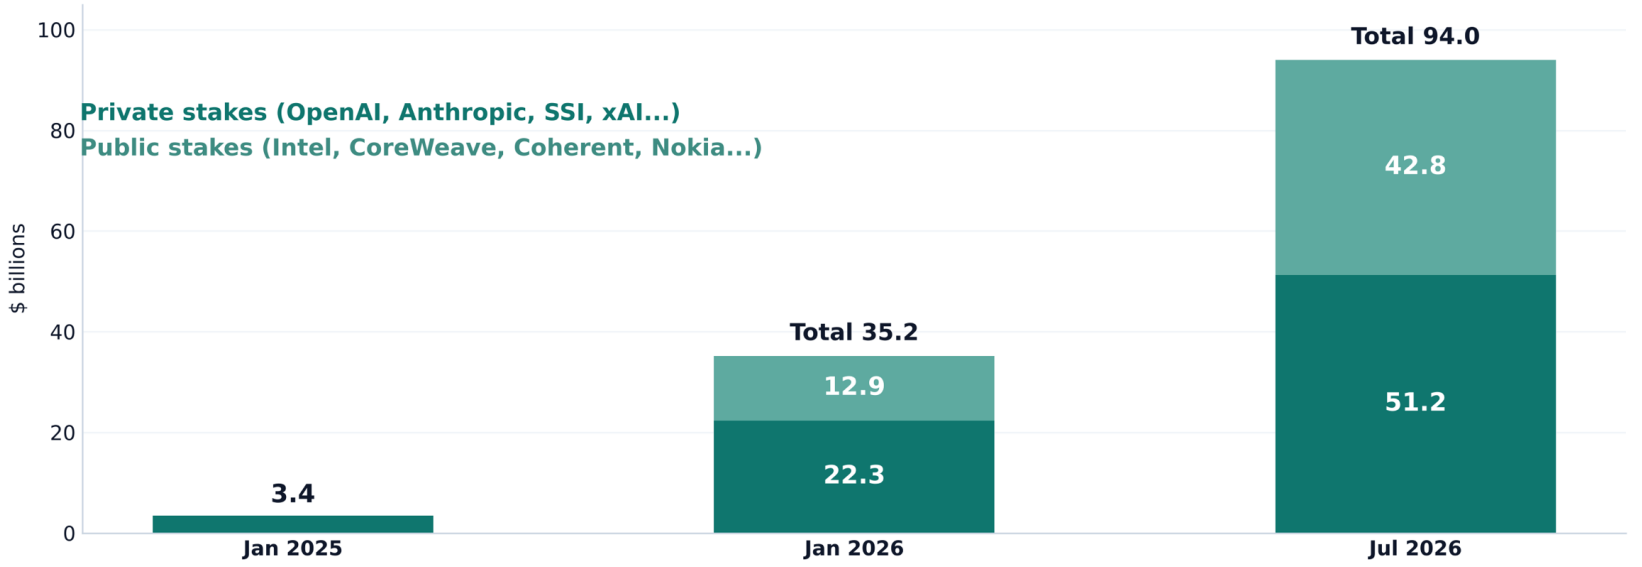
Huang on why NVIDIA finances AI, CNBC, Oct 2026

The venture arm: about $94B of stakes

ARC RESEARCH | AI UPDATE | NVIDIA THE INVESTOR

NVIDIA's investment book went from $3B to about $94B

Equity stakes on NVIDIA's balance sheet, $ billions: private (non-marketable) and public (marketable equity) holdings



Source: NVIDIA Form 10-K (Jan 25, 2026 and Jan 26, 2025) and Q2 FY27 results release (Jul 26, 2026), balance sheets. $ billions. Private = non-marketable equity securities. Public = marketable equity securities, shown separately from Jan 2026; in Jan 2025 they sat inside a combined marketable-securities line and are not shown. Total 94.0 = 51.2 + 42.8, ARC Research calculation. Excludes equity-method investments.

SO WHAT

Private stakes 15x in 18 months

NVentures: ~20 deals in 5 months

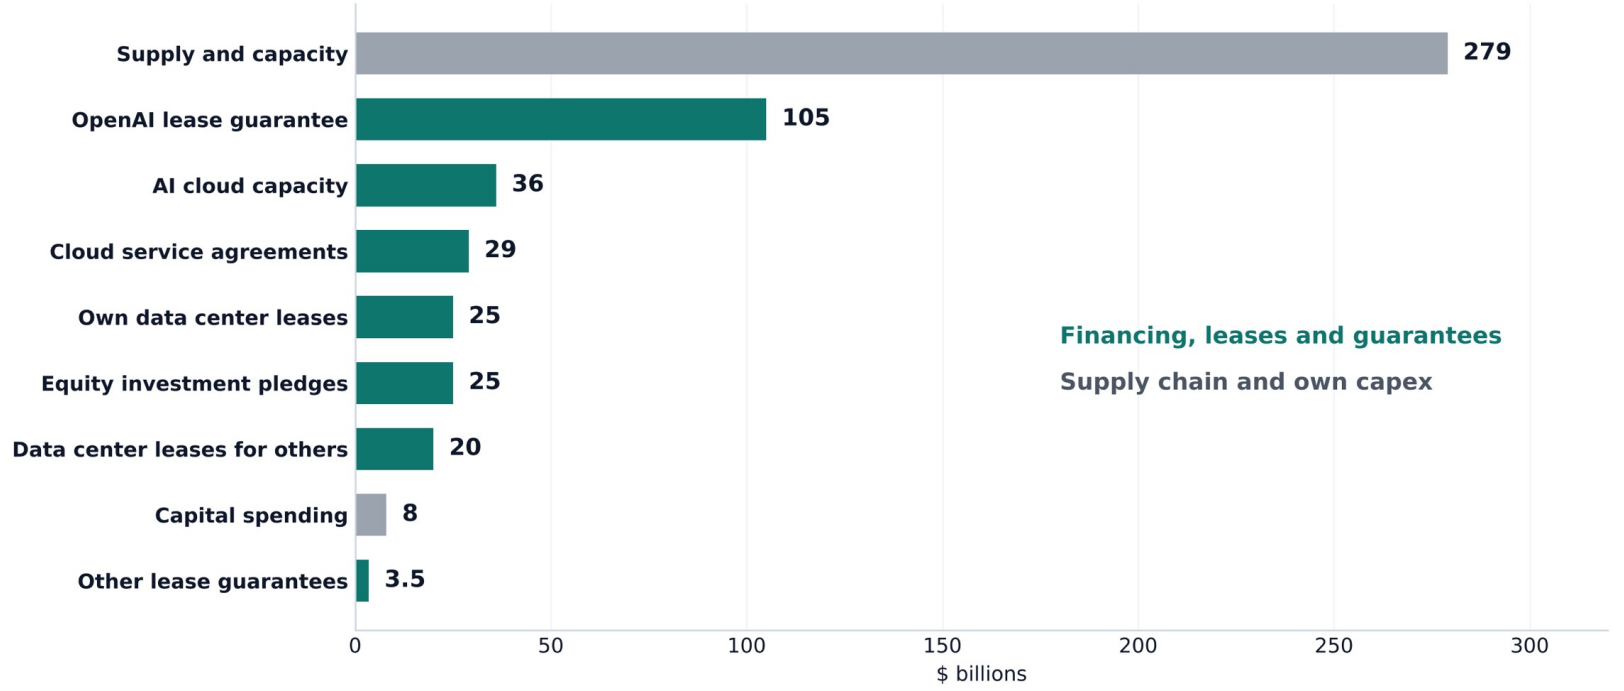
Public stakes fell ~$24B since June

The other side: $530B of commitments

ARC RESEARCH | AI UPDATE | NVIDIA COMMITMENTS AND GUARANTEES

About $530B of commitments sit off the balance sheet

Press-reported tally of the Jul 26, 2026 10-Q: supply orders dominate; $105B is a guarantee on OpenAI's leases



Source: Yahoo Finance (Nov 12, 2026) count from NVIDIA's Form 10-Q for the quarter ended Jul 26, 2026. $ billions.
Press-reported breakdown, not re-checked line by line against the filing. Items sum to $530.5B.

SO WHAT

Receivables $63.1B; 60 days to collect, was 45

Public stakes down about $24B since June

Kress signed the deals; now she sets the pace

Section 5 of 9

The interim CEO

Colette Kress: the CFO takes the wheel

COMING UP

29	Her career path
30	What to know about her now
31	The new board chair: Harvey Jones
32	Who is next: prediction markets
33	The favourite: Ian Buck
34	Her deal record: good and bad
35	How the stakes have done since June
36	13 years of NVIDIA’s numbers



Colette Kress: a career in tech finance



Colette Kress

Interim CEO · EVP and CFO

8 years	**Texas Instruments** · finance roles
13 years	**Microsoft** · CFO, Server and Tools
2010–13	**Cisco** · SVP and CFO, Business Technology and Operations Finance
Sept 2013	**NVIDIA** · EVP and CFO
Nov 2026	**NVIDIA · Interim CEO, and still CFO**

EDUCATION

BS Finance, Arizona · MBA, SMU Cox

MANDATE

Both jobs until a permanent CEO is named

STREET VIEW

A contender; less technical than Huang

Kress now: what to know

13 yrs
As NVIDIA’s CFO

~25 yrs
In tech finance before NVIDIA

2 jobs
Interim CEO and CFO at once

Nov 17
First earnings call as CEO

DAY ONE AND THE STREET

Day one
Gathers the executive team at HQ; tells staff they are NVIDIA’s greatest strength

The board
Full confidence; roadmap and customer commitments unchanged

The doubts
Less technical than Huang; markets favour Ian Buck for the permanent job

“She is absolutely the one that you need right now”
Jim Cramer, CNBC, Nov 12

Harvey Jones: NVIDIA’s first-ever board chair



Harvey C. Jones

Chair of the Board · director since 1993

1981–87	**Daisy Systems** · co-founder, CEO
1987–98	**Synopsys** · CEO to 1994, then executive chair
1993	**NVIDIA** · joins the board, its founding year
1997–2013	**Tensilica** · co-founder, chair; sold to Cadence
2004–	**Square Wave Ventures** · managing partner
Nov 2026	**NVIDIA · first-ever Chair of the Board**

BACKGROUND

Chip-design software CEO; BS Georgetown, MS MIT

ROLE NOW

Leads the board and the CEO search

WATCH

33 years beside Huang; independence questioned

Who is next? Markets favour Ian Buck

ARC RESEARCH | AI UPDATE | WHO IS NEXT

Prediction markets favour Ian Buck; Kress is second

Implied chance of being named NVIDIA's next permanent CEO, each dot a snapshot reported on Nov 12; interim CEOs do not count



Source: snapshots of Kalshi's "Who will be the next CEO of NVIDIA?" and Polymarket's matching market as reported on Nov 12, 2026 by CBS News, Kalshi News, RG.org (12:10 pm ET), FourWeekMBA (16:53 UTC) and Yahoo Finance. Prices move through the day and the markets are thin (Kalshi volume about $0.3M), so read these as a range, not a forecast. Kalshi resolves on an official NVIDIA announcement before Jan 1, 2029. Tim Cook's ~50% Polymarket print (CBS, Thursday morning) is an early outlier; Yahoo later put him near 6%. Not every outlet reported every candidate.

SO WHAT

Buck leads: ~32–42%

Kress second: ~12–24%

Thin markets; odds swing hourly

The favourite: Ian Buck, the engineer behind CUDA

Princeton	**BS computer science**
2000–04	**Stanford PhD** · led Brook, early GPU computing
2004	**Joins NVIDIA** after an earlier internship
2006	**Creates CUDA** · NVIDIA's software moat
2010s	**VP, Accelerated Computing** · all GPU-computing products
Today	**VP & GM, Hyperscale and HPC** · ~22 years at NVIDIA

MARKETS

32–42%

Odds of permanent CEO, Nov 12

WHY HIM

Runs sales to Microsoft, Google, AWS, Meta, OpenAI; advised the White House and Congress

QUESTIONS

Never a CEO; low public profile

Kress’s deal record: the good and the bad

GOOD CALLS

Mellanox, 2020

$7.13B buy; networking now $14.8B a quarter

Capital returns

$41.1B returned in FY26; dividend 25x

Caution on OpenAI

Reportedly pushed back on a $250B backstop

BAD OR MIXED

Arm, 2020–22

$40B deal blocked; $1.36B charge

H20 write-off, 2025

$4.5B charge after export rules changed

Circular financing

Stakes in customers; ~$24B mark-down since June

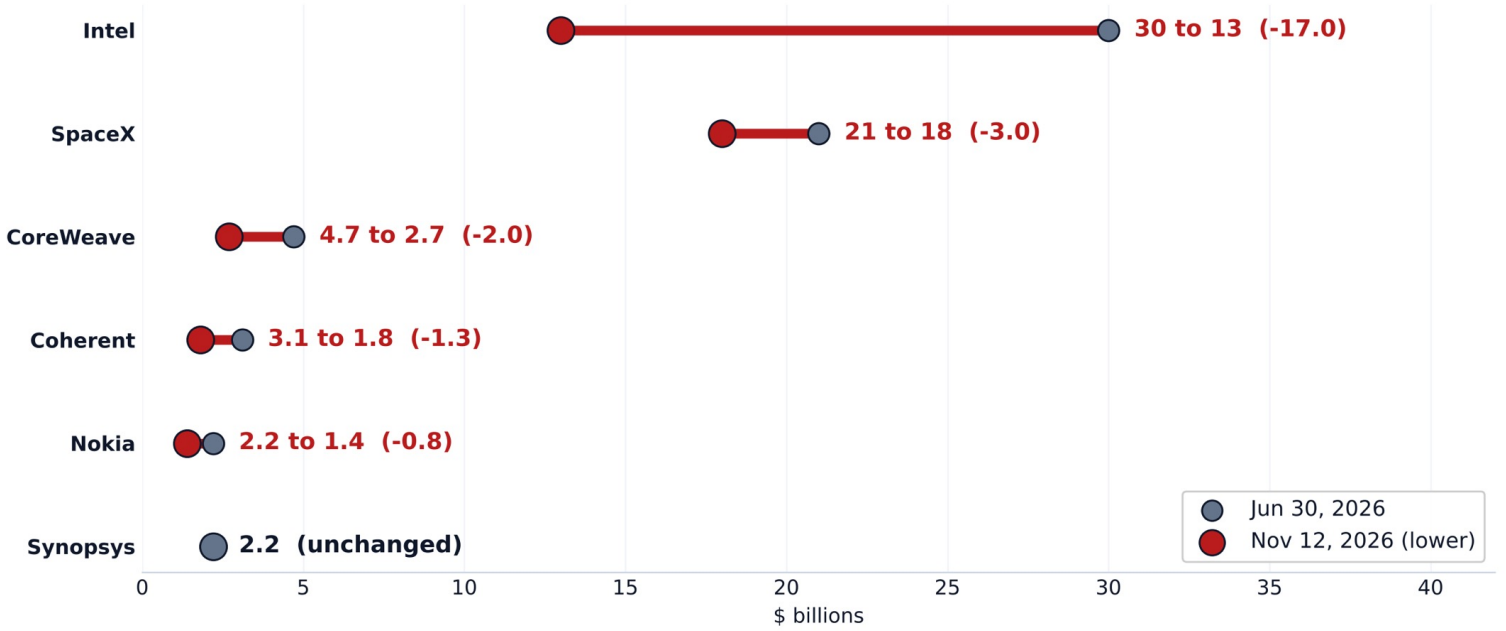
She signed off as CFO; many calls were Huang’s strategy. OpenAI pushback is press-reported.

Her investment book since June: about −$24B

ARC RESEARCH | AI UPDATE | THE CFO'S BOOK

NVIDIA's largest public stakes lost about $24B since June

Value of NVIDIA's largest public equity holdings, Jun 30 vs Nov 12, 2026, $ billions (approximate)



Source: Yahoo Finance, "Jensen Huang's exit puts Nvidia's $530 billion web of AI deals under the microscope" (Nov 12, 2026), chart of NVIDIA's largest public holdings.
Values are approximate as published by Yahoo Finance ("~"), not from NVIDIA filings; Nov 12 values are intraday. Change = Nov less Jun, ARC Research calculation.

SO WHAT

Intel stake: ~$30B to ~$13B

Losses hit reported earnings

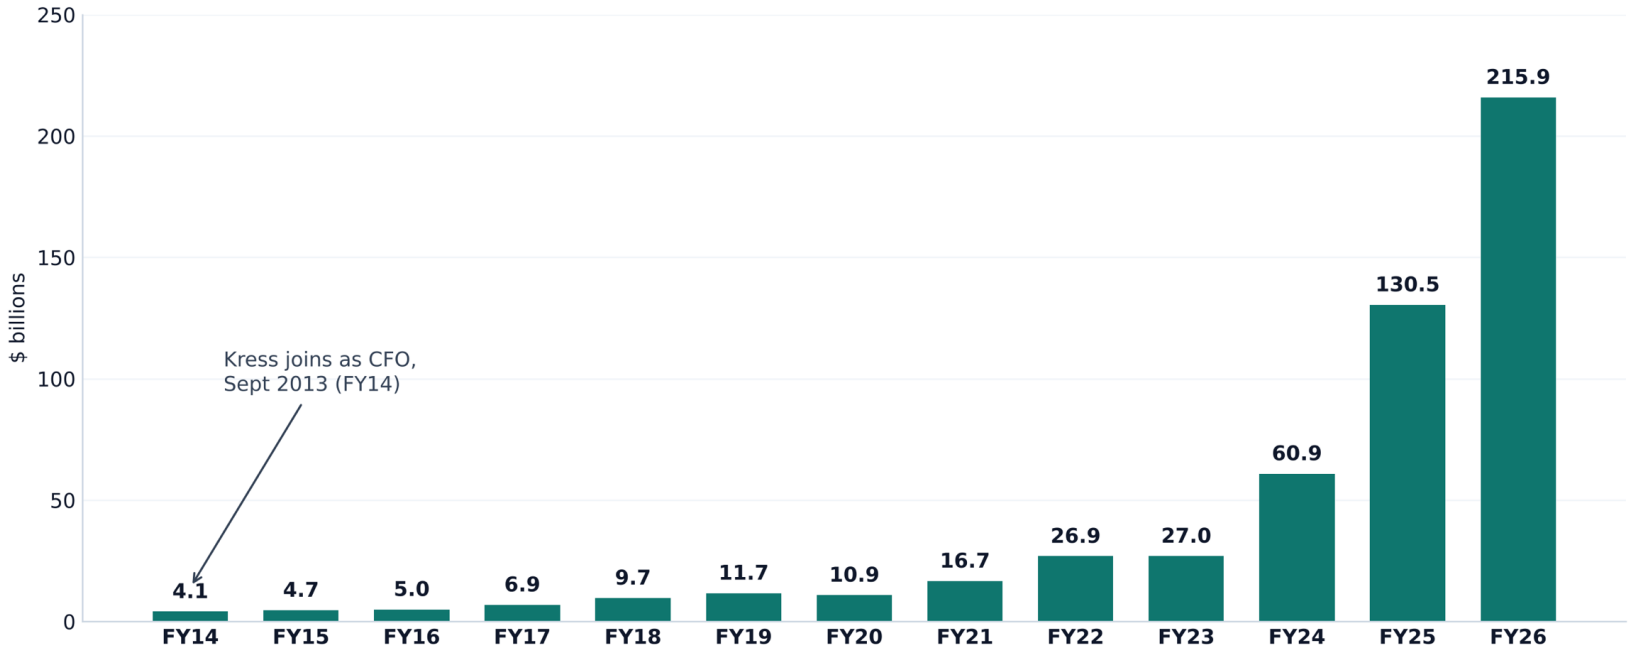
Keep, trim or add? Her call now

Kress’s 13 years as CFO: revenue up 52x

ARC RESEARCH | AI UPDATE | NVIDIA UNDER CFO COLETTE KRESS

NVIDIA annual revenue grew 52x in Kress's years as CFO

From $4.1B in fiscal 2014, the year she joined, to $215.9B in fiscal 2026



Source: NVIDIA annual reports (Form 10-K), fiscal years ending late January; FY26 from the Q4 FY26 results release. $ billions.
52x = 215.9 / 4.13, ARC Research calculation. FY14 to FY25 values are transcribed from the 10-Ks and were not re-fetched on Nov 12, 2026.

SO WHAT

$4.1B to $215.9B

Joined as CFO, Sept 2013

First call as CEO: Nov 17

Section 6 of 9

The Apple precedent

Apple when Steve Jobs died, and why Tim Cook took over

COMING UP

38 Why Apple had Tim Cook ready

39 Prepared vs unprepared: the handover

40 Product-market fit, so focus on supply

41 iPhone, iPad and Android growth

42 Demand: Apple then, NVIDIA now

43 Saturation vs technical demand

44 Competition then and now

45 Cook’s record: scale, not new categories

46 Two engines: supply chain vs new chips

47 Then and now, side by side



Why Apple had Tim Cook ready in 2011



Mar 1998	Jobs hires Cook from Compaq to fix operations
1998	Inventory cut from 30 days to 6
2004–09	Runs Apple during two Jobs medical leaves
Jan 2011	Acting chief for a third time
Aug 2011	**Jobs resigns for health; Cook named CEO on his recommendation**
Oct 2011	Steve Jobs dies, Oct 5

WHY COOK

Fixed chronic delays and supply problems

THE PARALLEL

Operator after visionary: Cook then, Kress now

THE DIFFERENCE

Cook was the planned successor; Kress is interim

Prepared vs unprepared: the handover



ACTING-CEO PRACTICE

3 vs 0

Cook’s stints before 2011 vs none at NVIDIA

WARNING TIME

7 mo vs 14 d

Last leave to handover vs report to resignation

SUCCESSOR

Plan vs search

Named by Jobs vs interim while the board looks

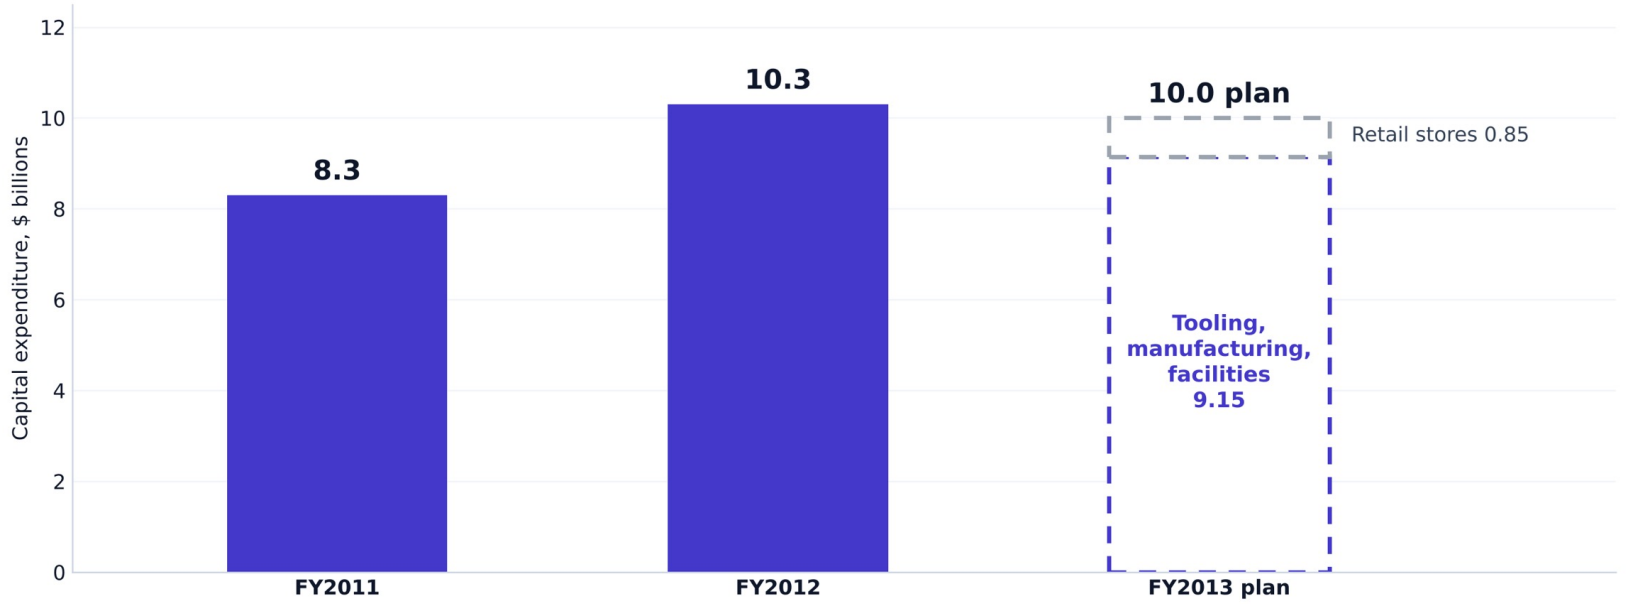
	Apple, 2011	**NVIDIA, 2026**
Successor	Cook, COO since 2007, 13 years at Apple	Kress, CFO, interim and still CFO
Succession plan	Jobs: “execute our succession plan”	No clear plan, per Fortune
Board	Jobs moves to chairman	First-ever chair named the same day
Exit	Planned; founder endorses successor	Sudden; under investigation

Product-market fit was won; the job was supply

ARC RESEARCH | AI UPDATE | APPLE AFTER PRODUCT-MARKET FIT

Demand was proven; Apple spent on supply

Capital spending by fiscal year: most of the 2013 plan was manufacturing tooling and equipment



Source: Apple fiscal 2012 Form 10-K, as summarised by Fortune (Nov 1, 2012). $ billions, fiscal years ending late September.
FY2013 is Apple's plan, not an actual (dashed): $9.15B for product tooling, manufacturing process equipment and corporate facilities, $0.85B for retail stores.

SO WHAT

Demand outran supply

Capex aimed at tooling and capacity

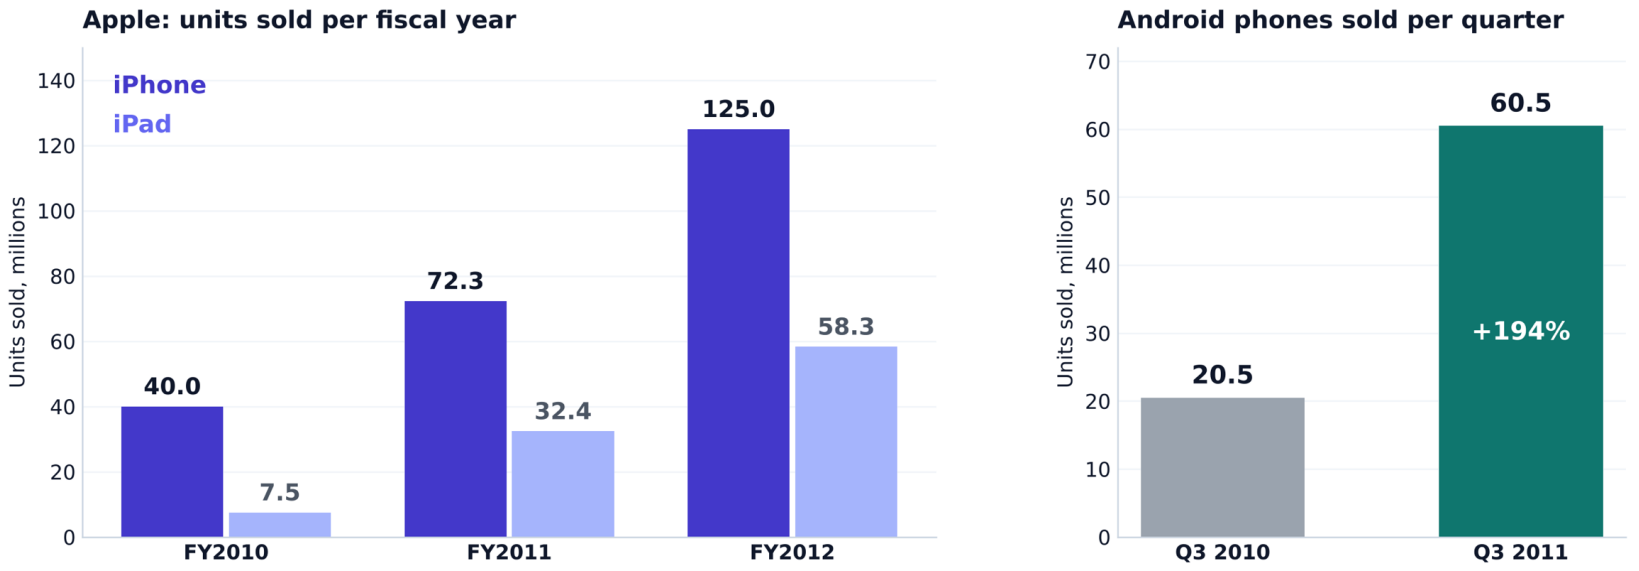
NVIDIA’s constraint is supply too

Both ecosystems were exploding

ARC RESEARCH | AI UPDATE | APPLE AND ANDROID AROUND 2011

Both ecosystems were exploding when Steve Jobs died

iPhone and iPad units tripled or more in two years; Android smartphone sales tripled in one



Source: Apple Form 10-K unit sales, fiscal years ending late September (FY2012 and FY2011 iPad as reported in the press; FY2010 and FY2011 iPhone transcribed, not re-fetched Nov 12, 2026). Android: Gartner smartphone sales to end users, 3Q10 and 3Q11 (press release, Nov 15, 2011). Millions of units.
Left panel is fiscal-year totals, right panel is one quarter; scales differ.

SO WHAT

iPhone: 40M to 125M a year

iPad: 7.5M to 58.3M a year

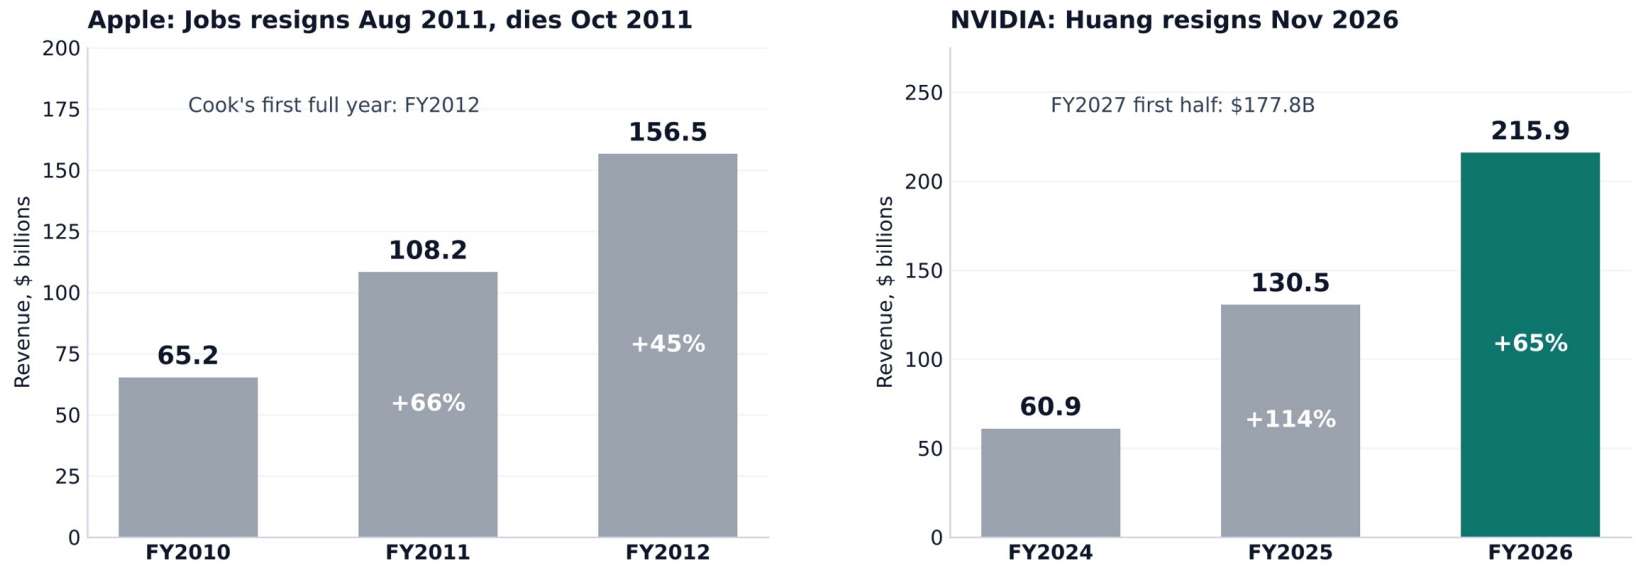
Android: 3x in one year

Both founders left at peak demand

ARC RESEARCH | AI UPDATE | FOUNDER EXITS: APPLE 2011 VS NVIDIA 2026

Both founders left at peak demand

Apple's revenue grew 66% in the year Steve Jobs died and 45% the year after; NVIDIA grew 65% last fiscal year



Source: Apple fiscal 2010-2012 results releases (fiscal years end late September; FY2011 $108.25B, FY2012 $156.5B). NVIDIA Form 10-K and Q4 FY26 release (fiscal years end late January). $ billions. Growth labels are year on year: Apple +66% as reported; +45% = 156.5 / 108.2 - 1; NVIDIA +114% = 130.5 / 60.9 - 1, +65% as reported. ARC Research calculations where derived. Different fiscal calendars; scales differ between panels.

SO WHAT

Apple: +66% the year Jobs died

+45% in Cook’s first full year

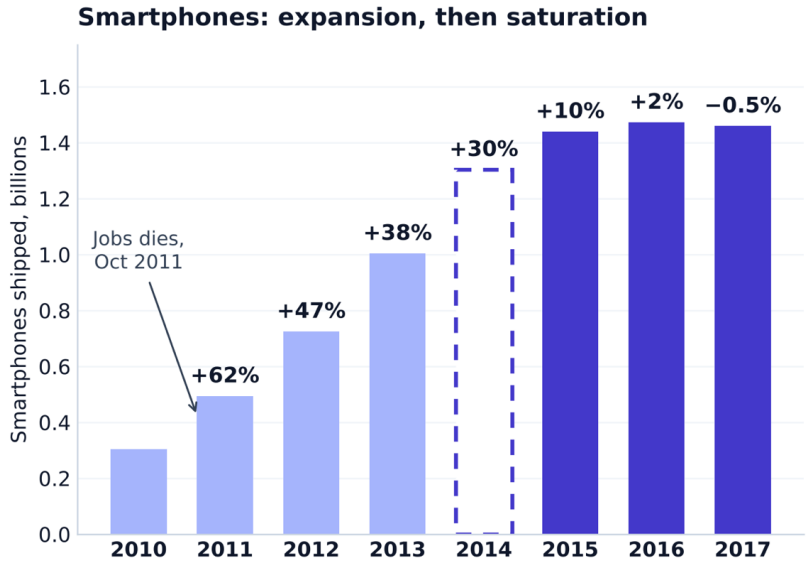
NVIDIA: +65%, then +106% last quarter

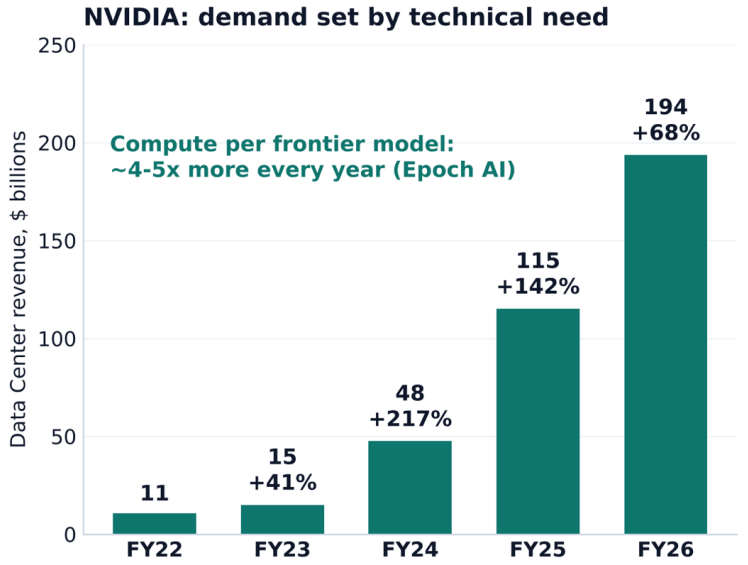
Smartphones saturated; AI demand is technical

ARC RESEARCH | AI UPDATE | TWO KINDS OF DEMAND

Smartphones ran out of new buyers; AI keeps needing more compute

Left: smartphone growth slowed to zero as most people owned one. Right: AI demand is set by compute per model, not by people





Source: IDC smartphone shipments as reported (2010 305M via Computerworld; 2011 494.4M, 2012 725.3M, 2013 1,004.2M; 2015 1.44B; 2016 1.47B, +2.3%; 2017 1.46B, -0.5%, IDC Feb 27, 2018 release). 2014 = 1,440 / 1.104, derived from IDC's +10.4% 2015 growth (ARC Research calculation, hollow bar). NVIDIA Data Center revenue: Form 10-K, fiscal years; FY26 = sum of four quarters. Compute trend: Epoch AI, training compute of frontier models grows about 4-5x per year (2010 to May 2024); a fitted trend, not NVIDIA data.

SO WHAT

Phones: growth 62% to zero by 2017

Buyers ran out, so focus on supply

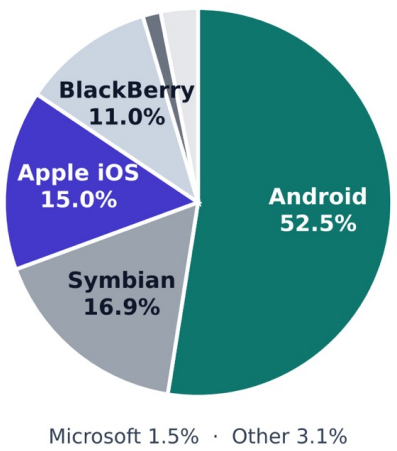
AI: each model needs ~4-5x more compute

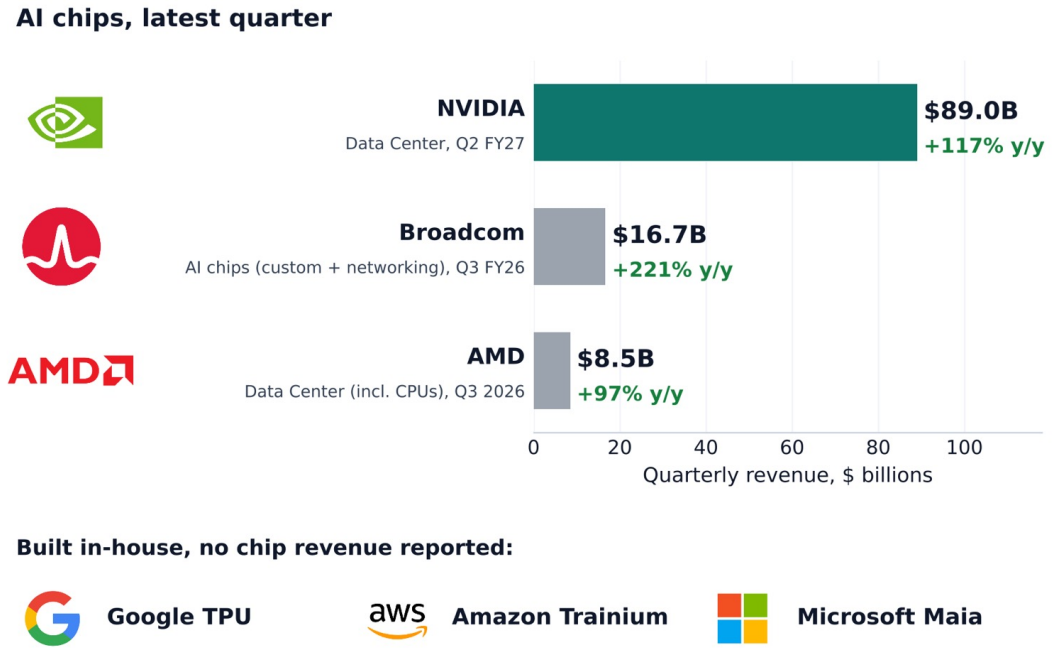
Competition then and now

ARC RESEARCH | AI UPDATE | COMPETITION THEN AND NOW

Then Android took half the market; now NVIDIA's rivals are smaller

Left: smartphone share when Jobs died. Right: latest reported AI chip revenue; rivals are a fraction of NVIDIA but grow faster





Built in-house, no chip revenue reported:

Google TPU · Amazon Trainium · Microsoft Maia

Source: left, Gartner smartphone sales to end users by OS, Q3 2011 (Nov 15, 2011 release, as reported by AFP and PhoneArena). Right, company filings: NVIDIA Data Center, Q2 FY27 (May-Jul 2026); Broadcom AI semiconductors, Q3 FY26 (to Aug 2, 2026; custom accelerators plus networking); AMD Data Center, Q3 2026 (to Sep 26, 2026; includes EPYC server CPUs). Segments are not like-for-like and quarters differ, so this is not a market-share pie. Growth is year on year as reported. Google, Amazon and Microsoft chips are built for their own clouds; no revenue is reported.

SO WHAT

2011: Android took 52.5%

Today: NVIDIA ~5x Broadcom, ~10x AMD

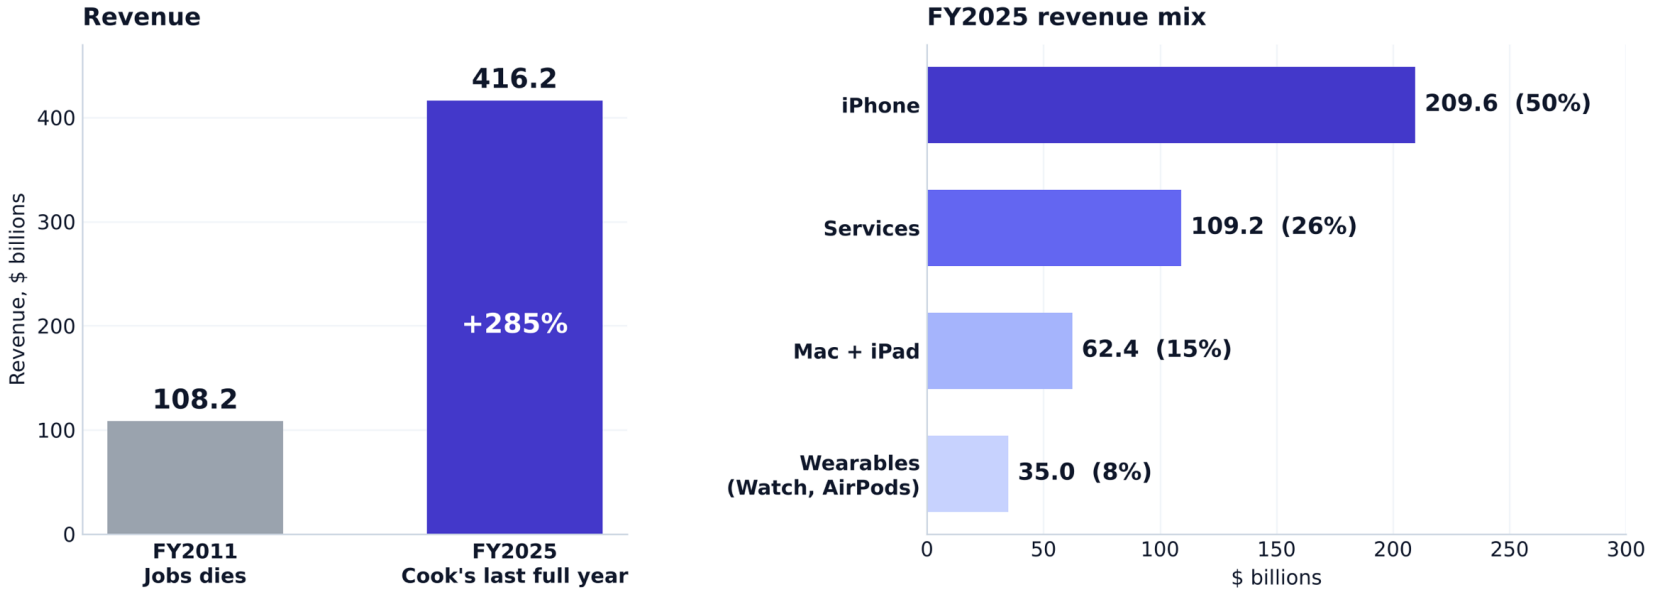
Rivals grow faster; clouds build their own

Cook scaled Apple; the next big product never came

ARC RESEARCH | AI UPDATE | THE COOK ERA AT APPLE

Cook nearly quadrupled revenue, but the iPhone is still half of it

Annual revenue, FY2011 vs FY2025, and the FY2025 mix: growth came from scale and Services, not a new iPhone-sized product



Source: Apple FY2011 results ($108.25B) and FY2025 results release (Oct 30, 2025: $416.2B; iPhone $209.6B and Services $109.2B as reported by Constellation Research). Wearables, home and accessories about $35B (Calcalist). Mac + iPad = 416.2 - 209.6 - 109.2 - 35 = 62.4, ARC Research calculation. +285% = 416.2 / 108.2 - 1, ours.
$ billions, fiscal years ending late September.

SO WHAT

Revenue ~4x; Services $109B

iPhone still 50% of revenue

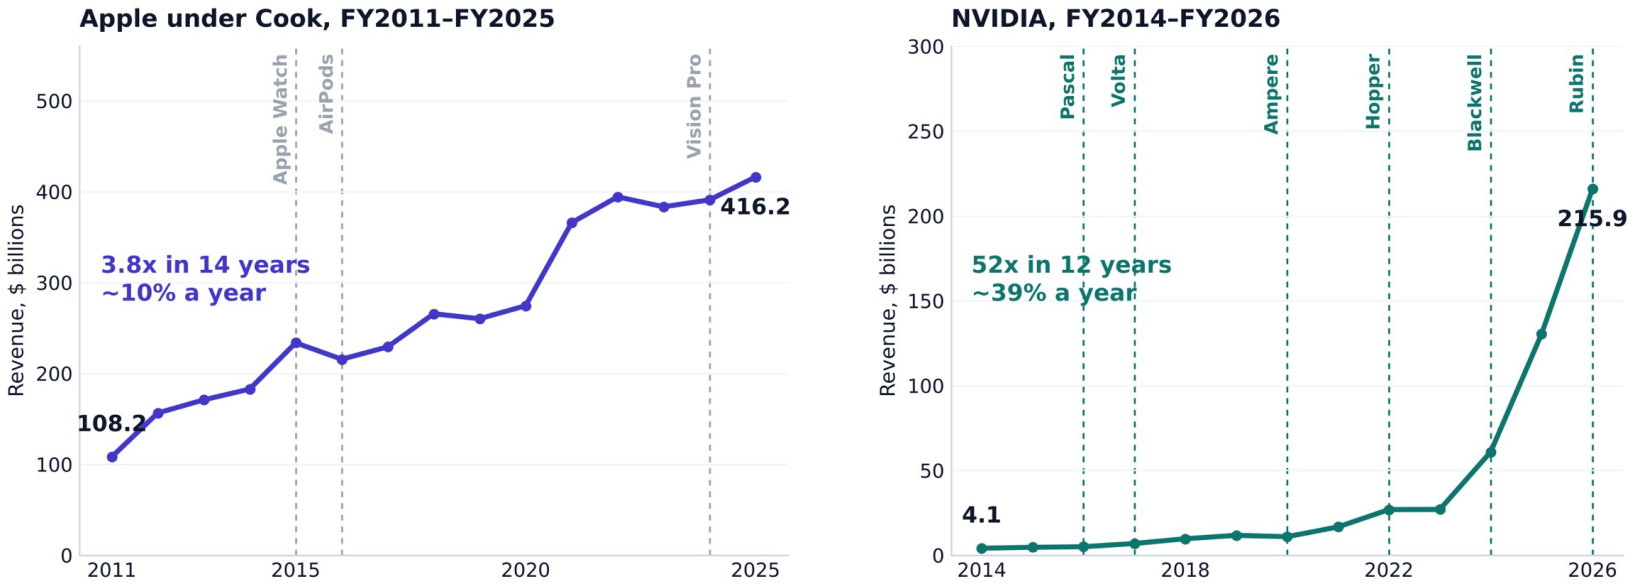
Critics: no iPhone-sized new product

Apple scaled a product; NVIDIA keeps reinventing it

ARC RESEARCH | AI UPDATE | TWO ENGINES OF GROWTH

Apple grew by scaling one product; NVIDIA by reinventing its chip

Annual revenue, $ billions. Markers: Apple's new product categories vs NVIDIA's new GPU architectures



Source: Apple and NVIDIA Form 10-K annual revenue (Apple fiscal years end late September, NVIDIA late January; NVIDIA FY2026 is the year ended Jan 2026). FY2011, FY2012, FY2025 Apple and FY2014, FY2024-26 NVIDIA checked against releases; other years transcribed from 10-Ks, not re-fetched Nov 12, 2026. Architecture years = launch year. CAGR = (end / start)^(1/years) - 1, ARC Research calculation. Scales differ between panels.

SO WHAT

Apple: ~10% a year from scale and supply

Few new categories after 2016

NVIDIA: ~39% a year from new chips

Then and now: Apple 2011 vs NVIDIA 2026

	Apple, 2011	NVIDIA, 2026
Exit	Jobs resigns Aug 2011; dies Oct 5	Huang resigns Nov 12 amid scandal
Demand	Revenue +66%, then +45%	Revenue +65%; +106% last quarter
Growth engine	Scale and supply chain	New chip every year
Rivals	Android at 52.5% share; Samsung No. 1	AMD Helios; Google, Microsoft, Amazon chips
Successor	Cook, operations chief, planned	Kress, CFO, interim; search under way
Balance sheet	$81.6B cash, no debt	$99.4B cash and securities; $33.4B debt
What followed	Revenue ~4x by FY2025; planned handover to Ternus, 2026	Open; next test is Nov 17

Read-across: demand looks like Apple in 2011, but Kress is interim and NVIDIA's rivals are its biggest customers

Section 7 of 9

The competition

Chips from Google, Microsoft, Amazon and AMD

COMING UP

49 Why chips must keep innovating

50 Four challengers

51 The technical race: chip specs

52 The revenue scoreboard



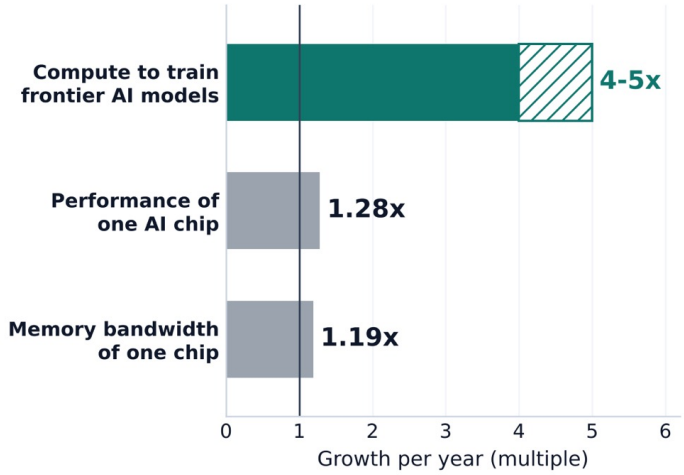
AI demands new chips every year

ARC RESEARCH | AI UPDATE | DEMAND FOR CHIP INNOVATION

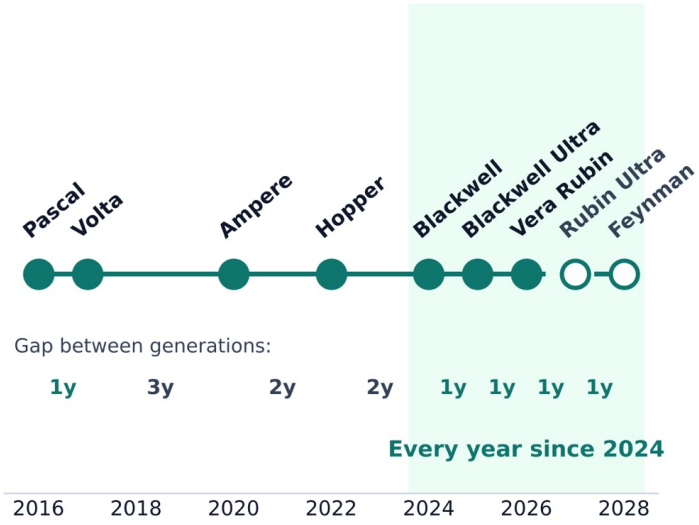
AI needs compute far faster than chips improve, so chips keep changing

Left: yearly growth in what AI wants vs what one chip delivers. Right: NVIDIA now ships a new data-center generation every year

Demand vs one chip, per year



NVIDIA data-center GPU generations



Over five years: ~1,800x more compute wanted vs ~3.4x from one chip. The gap is filled by more chips, faster links, new number formats, bigger memory stacks and new packaging: constant innovation.

Source: Epoch AI. Frontier training compute grows about 4-5x a year (2010 to May 2024). ML chip performance (16-bit FLOP/s) grows 28% a year, doubling every 2.8 years (Epoch ML hardware data, Nov 2024). Memory bandwidth doubles about every 4 years (Epoch, Nov 2023); 2^(1/4) = 1.19x a year, ARC Research calculation. Five-year multiples: 4.5^5 and 1.28^5, ARC Research calculation. Right: year each NVIDIA data-center GPU generation was introduced (NVIDIA announcements); 2027 and 2028 are roadmap targets from GTC 2026 (dashed).

SO WHAT

AI compute wanted: 4-5x a year

One chip improves only ~1.3x

NVIDIA now ships a generation a year

The challengers: four chips aimed at NVIDIA

GOOGLE

TPU 8 family

Four chips in one year

Two more launched Nov 9

Press-reported

MICROSOFT

Maia 300

Third generation

Unveiled in October

Press-reported

AMAZON

Trainium 3

3nm, 144GB of HBM3E

Runs most of Bedrock

Press-reported

AMD

MI450 · Helios

First rack revenue in Q3

Meta, Microsoft, OpenAI live

AMD Q3 2026 release

MARKET READ

AMD +13% on Nov 10; Alphabet +3% on Nov 12

THE CATCH

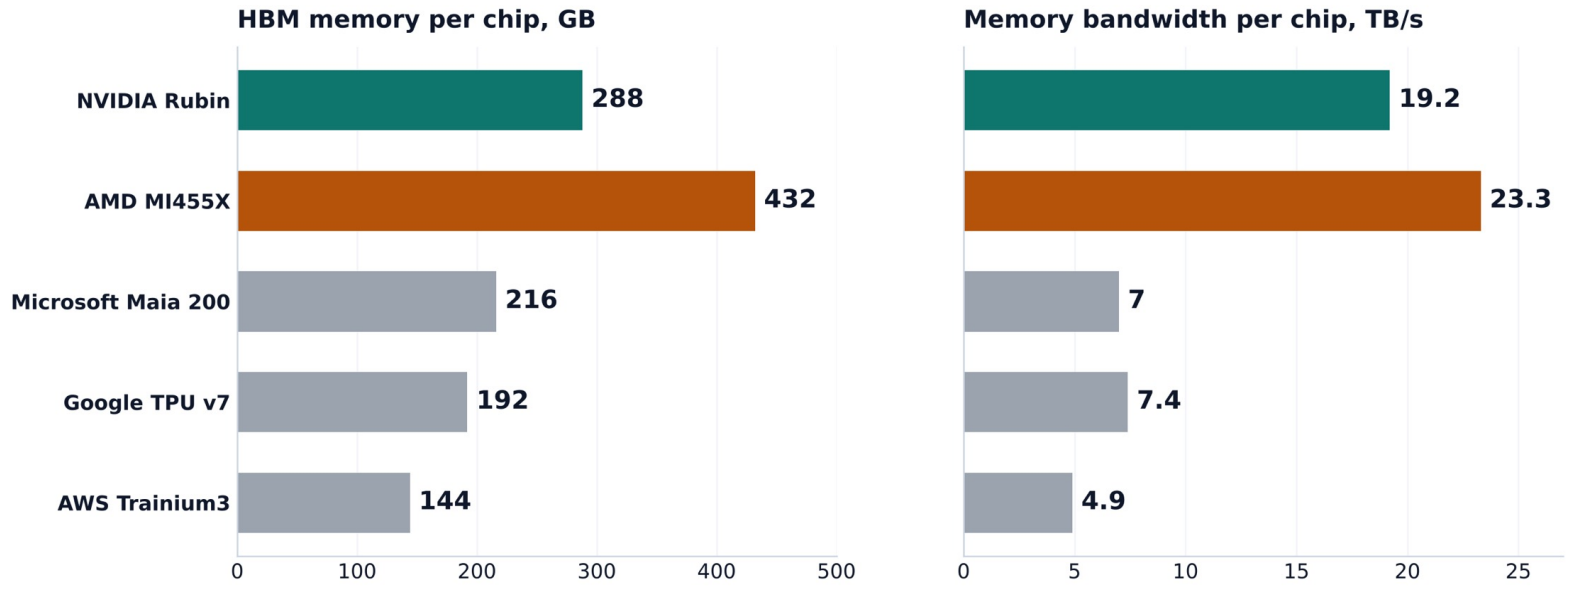
Three of the four are also top NVIDIA customers

The technical race: AMD leads Rubin on memory

ARC RESEARCH | AI UPDATE | THE TECHNICAL RACE

On paper, AMD now beats Rubin on memory; Google and Amazon trail

Published per-chip specs of each company's current flagship AI accelerator, as reported by the vendors



Source: NVIDIA Rubin spec page as revised (TweakTown, Wayback Machine; earlier listed 22 TB/s); AMD MI455X (AMD, via Tom's Hardware and StorageReview); Google TPU v7 Ironwood, AWS Trainium3, Microsoft Maia 200 (vendor specs as compiled by Techietory and Spheron, 2026). Published specifications, not benchmarks; not independently verified. Newer TPU 8 and Maia 300 specs were not available. FP4 compute: Rubin 50 PFLOPS, MI455X 40, Maia 200 ~10 (vendor figures).

SO WHAT

MI455X: 432GB vs Rubin’s 288GB

Rubin leads on FP4: 50 vs 40 PFLOPS

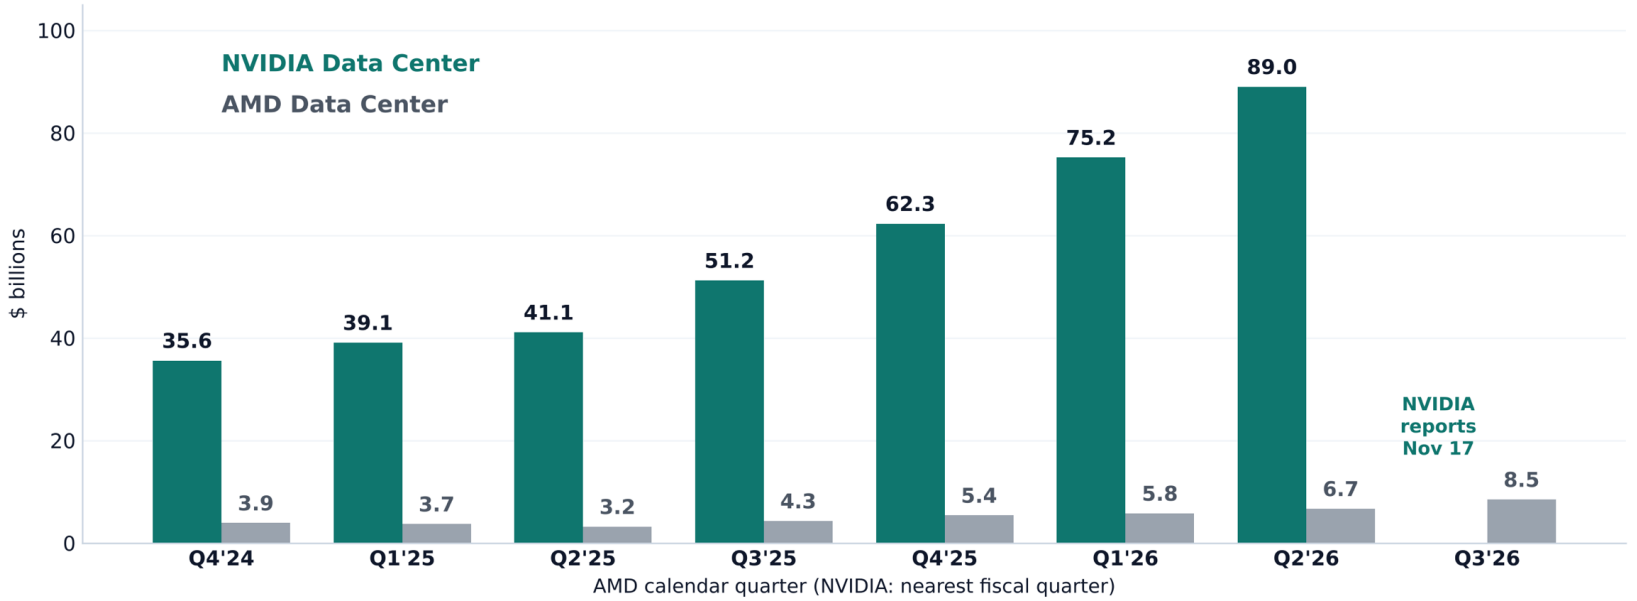
Specs, not benchmarks

The scoreboard: NVIDIA still about 10x AMD

ARC RESEARCH | AI UPDATE | THE MERCHANT-CHIP RACE

NVIDIA still earns about 10x AMD's data-center revenue

AMD's first Helios racks lifted it to $8.5B; NVIDIA's last reported quarter was $89.0B (next print Nov 17)



Source: NVIDIA and AMD quarterly results releases (SEC 8-K exhibits); quarters before 2026 as compiled in ARC Research AI Update, AMD vs NVIDIA (Nov 1, 2026).
$ billions. NVIDIA fiscal quarters end about a month after AMD's calendar quarters. AMD's segment includes EPYC server CPUs; NVIDIA's includes networking.
10x = 89.0 / 8.5, ARC Research calculation across non-identical quarters.

SO WHAT

NVIDIA $89.0B, +117% y/y

AMD $8.5B, +97% y/y

AMD: Helios to scale in Q4

Section 8 of 9

Shoes to fill

The supplier relationships Huang built in Taiwan and Korea

COMING UP

54 Taiwan and Korea ties

55 Jensen’s Taiwan trips

56 Who supplies the memory



Shoes to fill: Huang’s supplier ties

TAIWAN

TSMC

Makes the chips. Frequent Huang visits to TSMC, Foxconn and Wistron

KOREA

SK hynix

Lead HBM supplier. Ties to SK chairman Chey Tae-won

KOREA

Samsung

Second HBM4 source. Ties to chairman Lee Jae-yong

“the new CEO may not be as aggressive as Jensen”

Taiwan supplier executive, to Reuters, Nov 12

THE RISK

Slower capacity pre-booking and spec upgrades

THE BUFFER

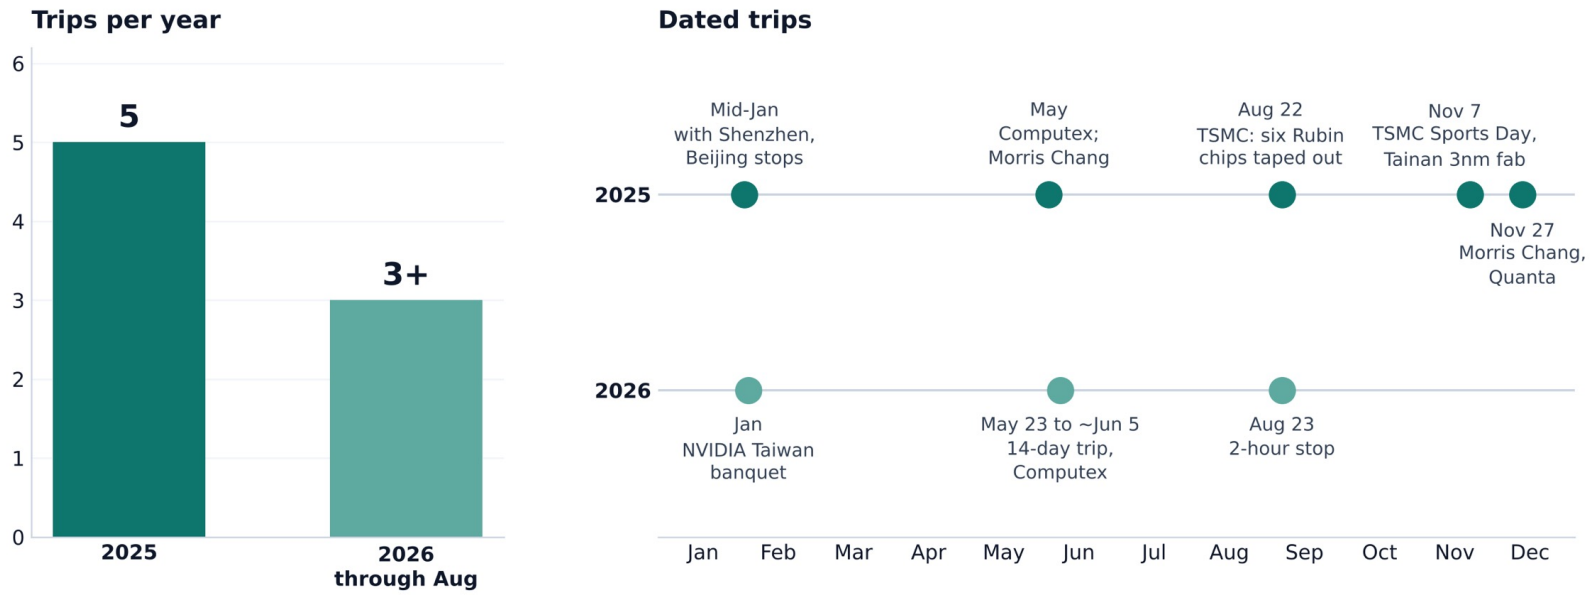
Long-term HBM deals; supply still tight

Jensen in Taiwan: five trips, weeks on the ground

ARC RESEARCH | AI UPDATE | JENSEN IN TAIWAN

Five Taiwan trips in 2025; one 2026 trip lasted 14 days

Huang's press-reported visits to Taiwan, mostly to TSMC, suppliers and Morris Chang



Source: Taipei Times and Focus Taiwan (Aug 22-23, 2025: third trip of 2025); Baidu Baike summary of Taiwanese press (Nov 7, 2025: fourth); United Daily News via Asia Business Daily and Digitimes (Nov 27, 2025: fifth); Yicai (Jan 2025: Taiwan among China stops); Baidu Baike (May 2025 with Morris Chang); Taiwan News (Aug 2026: third of 2026); Focus Taiwan (May 23, 2026 arrival, 14-day trip). Press counts, not company data; trips not reported in the press would be missing, so counts may be low.

SO WHAT

5 trips in 2025; 3+ by Aug 2026

May 2026 trip lasted 14 days

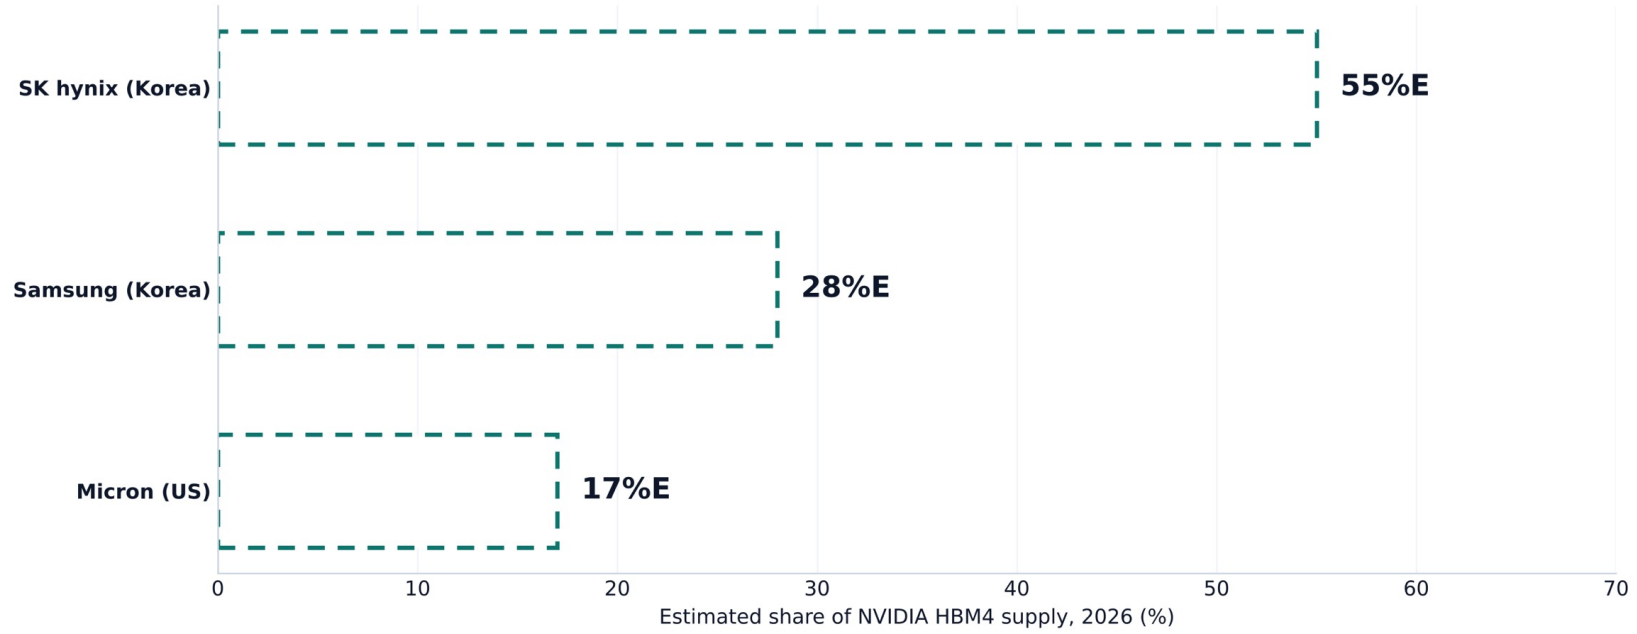
Press counts; may be low

Memory: Korea holds most of NVIDIA’s HBM4

ARC RESEARCH | AI UPDATE | NVIDIA'S MEMORY SUPPLIERS

Korea supplies most of NVIDIA's HBM4, led by SK hynix

Analyst estimate of NVIDIA's 2026 HBM4 supply by vendor; SK hynix says capacity is sold out through 2029



Source: Daishin Securities estimate, as reported by Korea JoongAng Daily (Nov 12, 2026). Share of NVIDIA's sixth-generation HBM4 supply in 2026, %.
Hollow bars = analyst estimate, not company-reported. Sold-out comment: SK hynix Q3 2026 earnings call, per the same report.

SO WHAT

SK hynix about 55%, Samsung 28% (estimates)

SK hynix sold out through 2029

Kress inherits these relationships

Section 9 of 9

ARC Research view

What we think happens next: a view, not advice

COMING UP

58	Not a Tim Cook moment
59	NVIDIA stays ahead; what the edge is worth
60	Short term fine; long term the question
61	Demand will be met; rivals catch up
62	The likely end of circular funding
63	Huang’s open-source push may stall
64	Unconventional bets meet a CFO
65	Buyers diversify; Cerebras grows
66	Labs back tighter export rules
67	The ARC Research view in nine points



ARC view: this is not a Tim Cook moment

Apple, 2011

SUCCESSOR	**Planned:** Cook, picked by Jobs
TRACK RECORD	Ran Apple 3 times during Jobs's leaves
THE JOB	Product fit found; **scale the supply chain**
LEADER NEEDED	An operator: Cook's strength

NVIDIA, 2026

SUCCESSOR	**None:** CFO as interim; search open
TRACK RECORD	No heir groomed in 33 years
THE JOB	**Better chips every year**; grow CUDA ecosystem
LEADER NEEDED	A technical, product-led builder

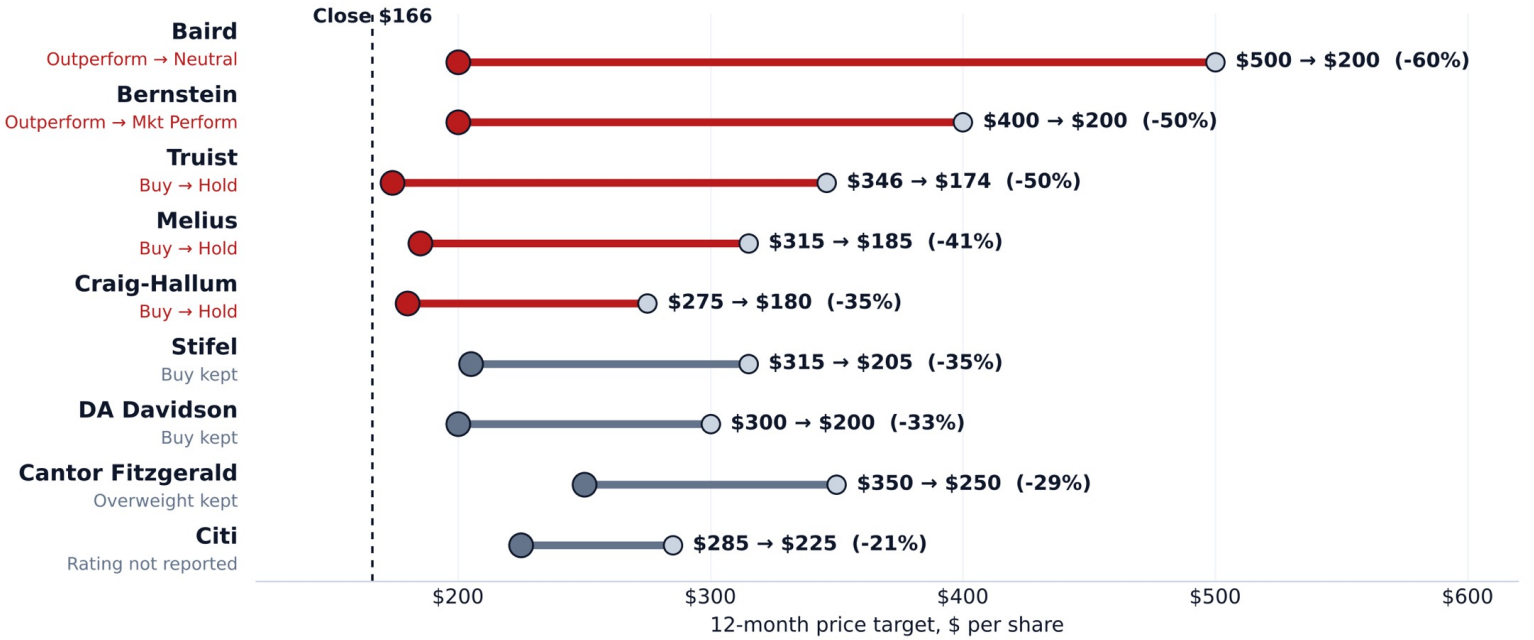
Our read: Apple needed an operator to scale a finished product; NVIDIA needs a builder to keep out-innovating

ARC view: NVIDIA stays ahead; what is the edge worth?

ARC RESEARCH | AI UPDATE | ANALYST ACTIONS

Wall Street cut targets by up to 60% in two days

Price targets before and after Huang's exit, Nov 11-12; red = rating also downgraded. The close on Nov 12 was $166.19



Source: CNBC, Motley Fool, MLQ.ai, Finbold and TokenPost reports of analyst notes, Nov 11-12, 2026 (press-reported; not every note was seen). Baird's cut is Nov 11; the rest Nov 12. Downgraded with targets removed: Oppenheimer (to Perform) and Evercore ISI's Mark Lipacis (to In Line). William Blair downgraded to Market Perform. Citi's rating change, if any, was not reported. Close: $166.19 on Nov 12 (Motley Fool). Cut % = new / old - 1, ARC Research calculation.

SO WHAT

Not going away: still the best GPUs

Question: how long the lead lasts

Street now pays less for it

ARC view: fine in the short term; the long term is the question

Short term: fine NEXT 12–18 MONTHS

DEMAND	Supply-constrained; Q3 guide $108B
OUTLOOK	~70% growth guided for FY28
PRODUCT	Vera Rubin ramping now
SUPPLY	$279B of supply locked in

Long term: the question 2028 AND BEYOND

LEADERSHIP	No permanent CEO; no heir
ROADMAP	Rubin Ultra, Feynman without Huang
RIVALS	Buyers diversify; clouds build chips
PRICE	What is the edge worth? Targets cut

Our read: the next year is largely booked; the risk is a slow erosion of the lead after 2027

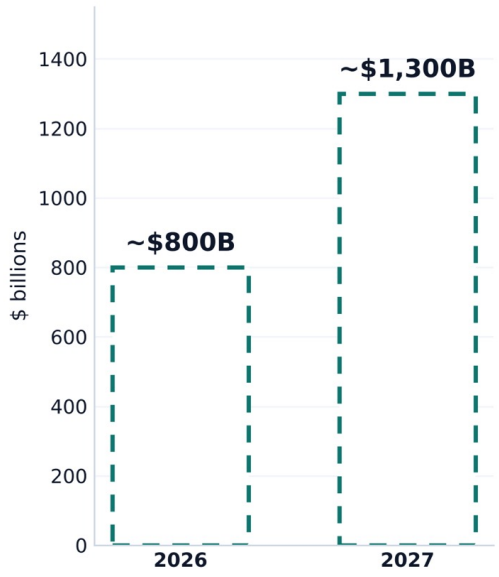
ARC view: the demand will be met by someone

ARC RESEARCH | AI UPDATE | ARC RESEARCH VIEW

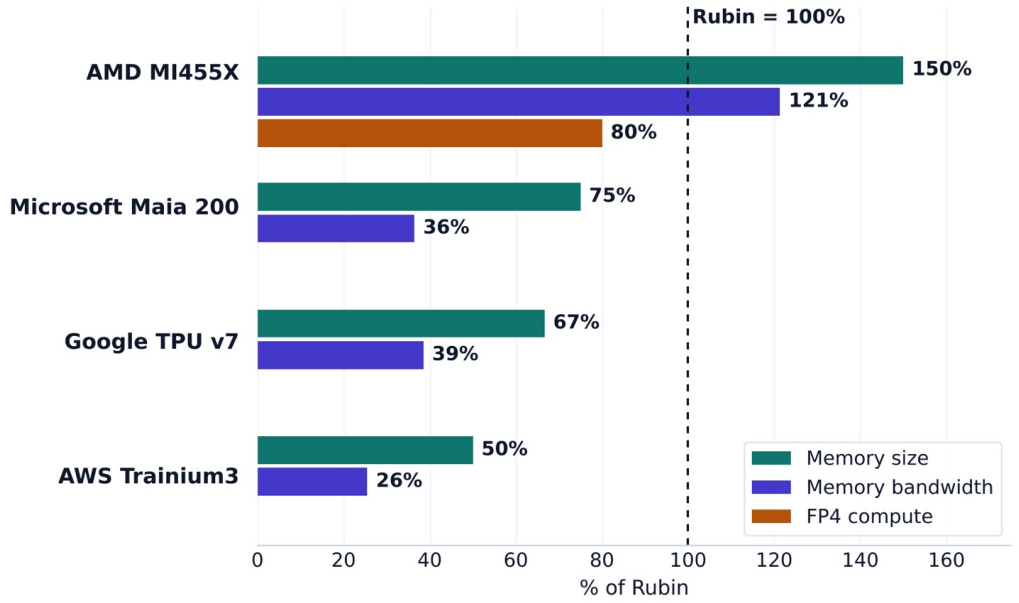
The demand will be met by someone, and rivals are closing the gap

Left: what the biggest buyers expect to spend. Right: rival chips' published specs as a share of NVIDIA's Rubin (100% = parity)

Top-5 hyperscaler capex



Rival chips vs NVIDIA Rubin, published specs



Source: left, NVIDIA Q2 FY27 earnings call (Aug 26, 2026): top-5 hyperscaler capex expected to reach nearly $800B in 2026 and $1.3T in 2027 (expectations, dashed). Right: vendor specifications as reported: Rubin 288 GB HBM4, 19.2 TB/s (revised from 22), ~50 PFLOPS FP4; AMD MI455X 432 GB, 23.3 TB/s, 40 PFLOPS FP4; Google TPU v7 192 GB, ~7.4 TB/s; Microsoft Maia 200 216 GB, ~7 TB/s; AWS Trainium3 144 GB, 4.9 TB/s (TweakTown, Tom's Hardware, StorageReview, Network World). Ratios = rival / Rubin, ARC calculation. Specs, not benchmarks; cloud FP4 not comparable.

SO WHAT

~$1.3T of capex wants chips in 2027

If not NVIDIA, someone supplies it

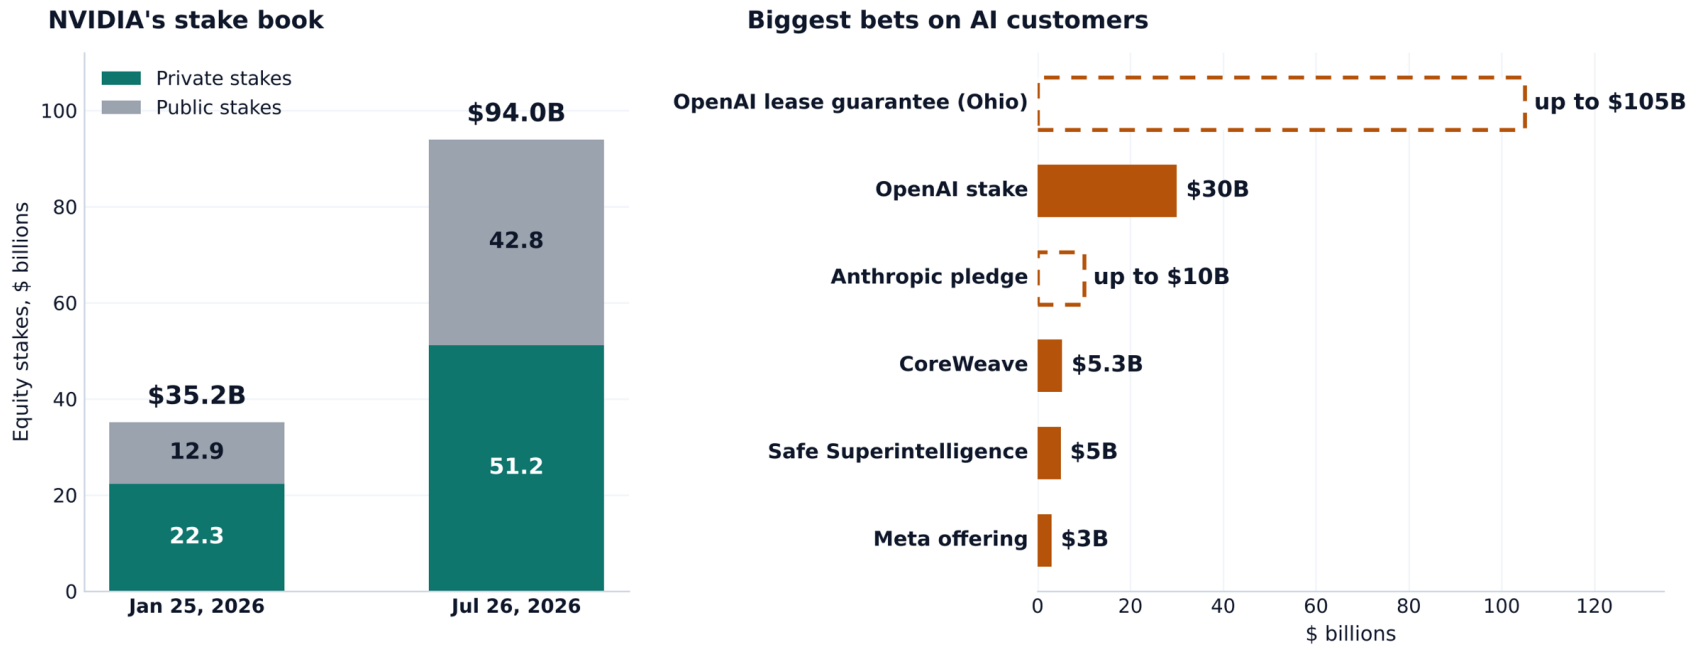
AMD already leads Rubin on memory

ARC view: the likely end of circular funding

ARC RESEARCH | AI UPDATE | ARC RESEARCH VIEW

Huang built a $94B stake book in customers; expect it to stop growing

Left: NVIDIA's equity stakes nearly tripled in six months. Right: the biggest bets on AI companies that also buy its chips



Source: left, NVIDIA balance sheet (Q4 FY26 and Q2 FY27 releases): marketable equity securities (public) and non-marketable equity securities (private), $B.
Right: Yahoo Finance tally of NVIDIA deals (Nov 12, 2026; press-reported). Solid = invested; dashed = pledged or guaranteed (Anthropic pledge up to $10B; OpenAI Ohio lease guarantee up to $105B).
CoreWeave = $3.3B stake plus $2B more. 42.8 + 51.2 = 94.0, ARC Research calculation. Values of public stakes move with share prices.

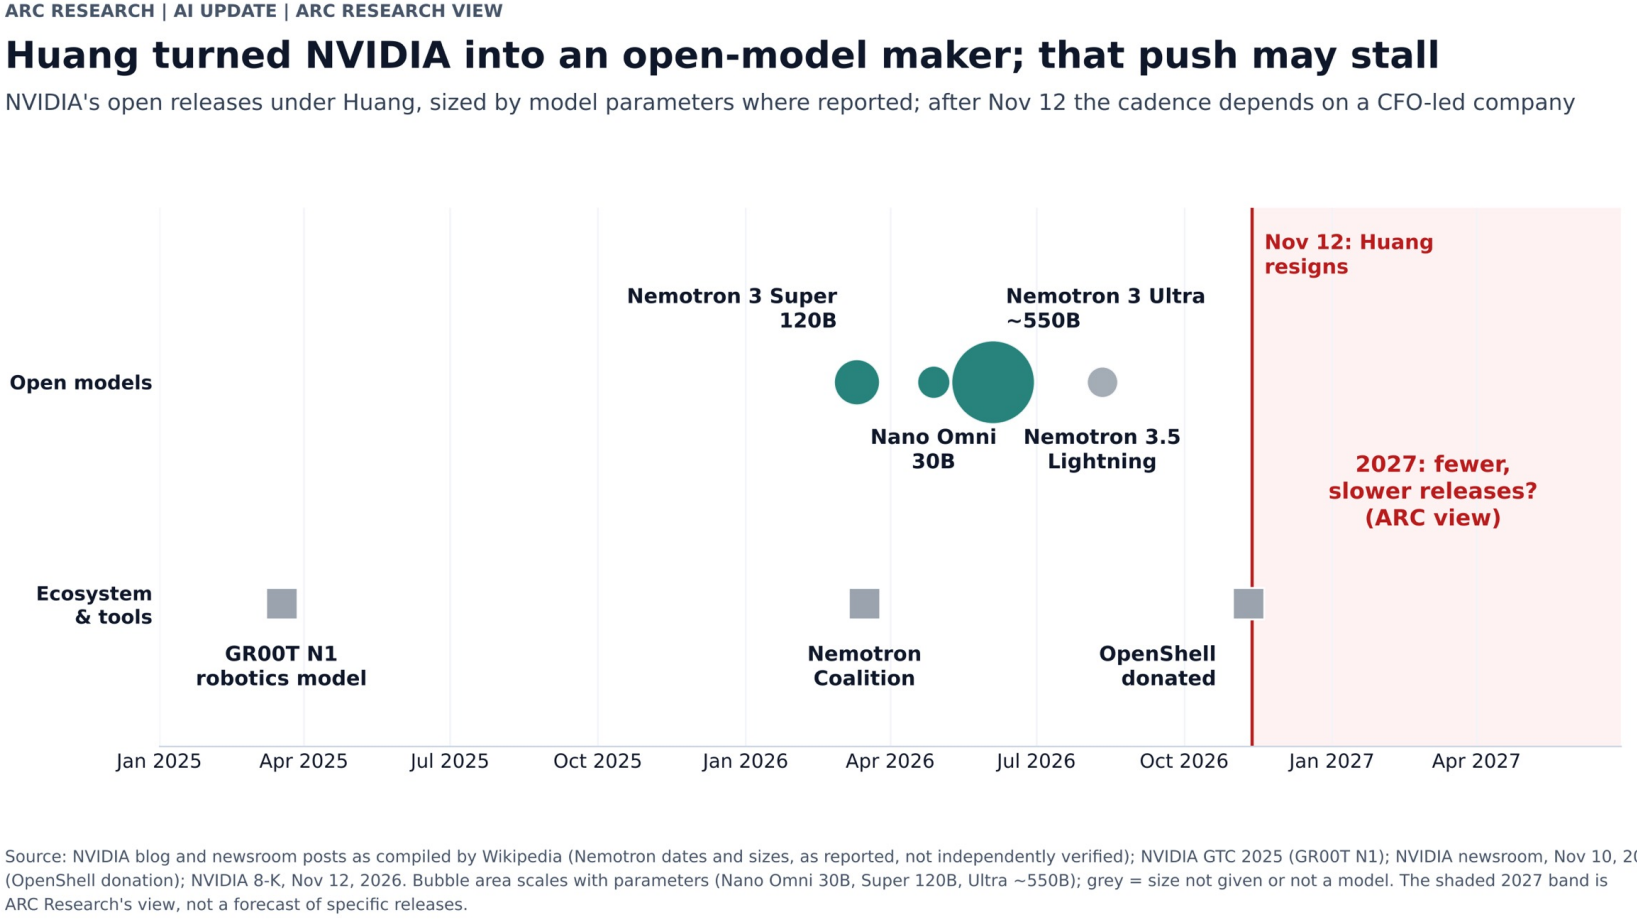
SO WHAT

Huang funded customers to buy chips

A CFO-CEO puts returns first

Expect few new stakes or guarantees

ARC view: Huang’s open-source push may stall



SO WHAT

Huang made NVIDIA an open-model maker

4 Nemotron releases in 5 months

Costly, no direct revenue: at risk

ARC view: unconventional bets meet a conventional CFO

The bet	Huang's logic	A traditional CFO's lens	Our expectation
Stakes in AI labs	Fund customers so they buy chips	Balance-sheet risk; looks circular	Few new stakes
Lease guarantees	Unblock data centers faster	Up to $105B contingent liability	No new big guarantees
Open-source models	Keep open AI running on NVIDIA	Costly; no direct revenue	Slower, fewer releases
Long-shot platforms	Robotics, quantum, 6G bets	Long payback; hard to measure	Trimmed to the core

HUANG

Founder; visionary bets a decade early

KRESS

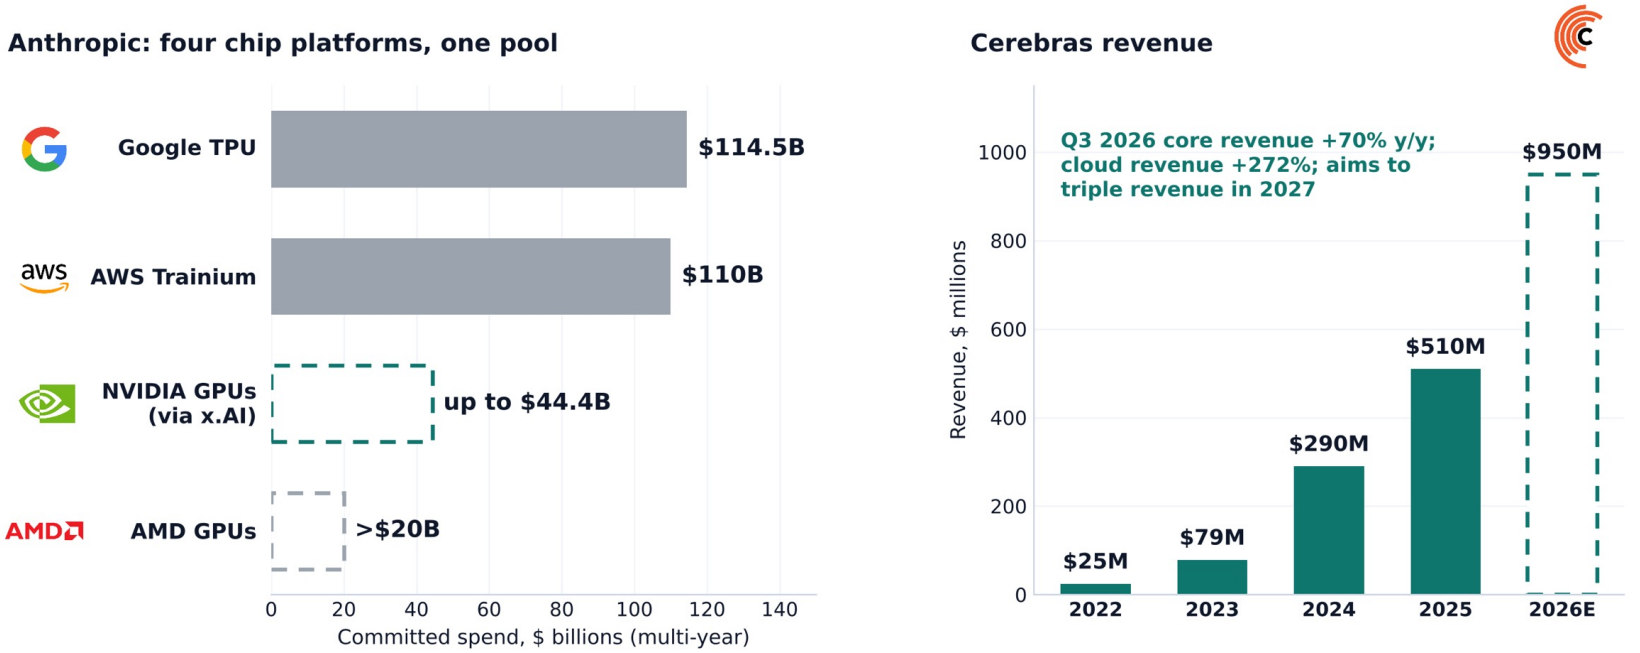
Career CFO: TI, Microsoft, Cisco, NVIDIA

ARC view: buyers diversify; challengers grow

ARC RESEARCH | AI UPDATE | ARC RESEARCH VIEW

Buyers are spreading bets across chips, and challengers grow

Left: Anthropic's disclosed compute commitments by chip platform. Right: Cerebras, a wafer-scale challenger, is set to nearly double revenue



Source: left, Anthropic S-1 as reported by CNBC (Oct 26, 2026): $110B on AWS Trainium through 2036, $114.5B on Google TPUs through 2033, could buy more than $20B of AMD chips (dashed); NVIDIA GPUs via x.AI up to $44.4B, May 2026 to May 2029, cancellable on 90 days' notice (Anthropic DRS/A; dashed). Other NVIDIA capacity (Microsoft, Nscale, Lambda) is not split out. Durations differ, so these are not annual spending shares. Right: Cerebras prospectus (424B4) for 2022-2025 GAAP revenue; 2026 = core revenue guidance $940-960M, midpoint shown dashed (Q3 2026 release, Nov 12).

SO WHAT

Anthropic runs on 4 chip platforms

NVIDIA is one supplier of several

Cerebras: ~$950M in 2026; 3x planned

ARC view: the labs back tighter chip export rules



Anthropic: backs all three
Publicly; quit a trade group over it



OpenAI: now backs them
Told lawmakers this week, after going quiet in October

THE BIPARTISAN BILLS (LIKELY RIDING ON THE DEFENSE BILL)

AI OVERWATCH Act

Congress reviews chip export licences, like arms sales

Backed by Anthropic and OpenAI

Chip Security Act

Track exported AI chips to stop smuggling

Backed by Anthropic and OpenAI

MATCH Act

Push allies to curb chip tools to China

Backed by Anthropic and OpenAI

Our read: fewer NVIDIA sales to China; US labs keep first call on supply. Senate leader Thune: “high probability”

ARC Research view (a view, not advice)

01

Founder era over

33 years, one CEO, a $4T company

02

Not a Tim Cook moment

Apple had an heir to scale supply; NVIDIA has none and must out-innovate

03

Fortress finances

$66B net cash; $530B committed

04

Less circular funding

A traditional CFO cuts stakes in customers

05

Fewer open-source bets

A CFO-CEO trims Huang’s long shots

06

Rivals gain demand

AMD, TPU, Trainium, Maia win orders

07

TSMC capacity reshuffles

New chip companies fill the slots

08

Tighter chip rules

China sales curbed; US labs first

09

What to watch

Supplier ties, earnings Nov 17

Sources (1 of 2)

NVIDIA 8-K: CEO transition, Nov 12, 2026

NVIDIA Q2 FY27 results release (8-K), Aug 26, 2026

NVIDIA Q1 FY27 and Q4 FY26 results releases (8-K)

NVIDIA 10-K annual reports, FY14–FY25 (not re-fetched)

NVIDIA: Kress appointed CFO, Sept 2013

Korea Economic Daily: Huang's influence in Washington

Wired: NVIDIA becomes a major model maker with Nemotron 3

NVIDIA newsroom: OpenShell donation, Nov 10, 2026

TweakTown: NVIDIA revises Vera Rubin bandwidth

Tom's Hardware: AMD Instinct MI455X and Helios

Focus Taiwan: Huang ends 14-day Taiwan trip, Jun 2026

Fox Business: Warren, Coons demand Nvidia records

NBC News: Megaspeed pledges to cooperate

CNBC live updates: Huang to testify, Nov 12

NVIDIA 8-K: Arm termination, Feb 8, 2022

The Register: H20 $4.5B charge, May 2025

AI Weekly: Kress pushed back on OpenAI backstop

Taipei Times: Huang visits TSMC, Aug 23, 2025

Investing.com: NVDA historical prices, Nov 12

Yahoo Finance: $530B web of AI deals, Nov 12

NBC News: stock falls after Huang resigns, Nov 12

Motley Fool: timeline and Citi target, Nov 12

Forbes: timeline of the week, Nov 12

Fortune: Huang resigns, Kress interim, Nov 11

CNBC: Who is Colette Kress?, Nov 12

CNBC: NVIDIA as AI financier, Oct 24 and Aug 18

AMD Q3 2026 results release (8-K), Nov 3, 2026

Apple FY2011 results, via MacRumors, Oct 2011

Fortune: Apple's fiscal 2012 annual report, Nov 2012

Gartner Q3 2011 smartphone sales, via AFP

Sources (2 of 2)

Apple FY2025 results release, Oct 30, 2025

Adweek: Tim Cook grew Apple by reducing its ambition

AppleInsider: Tim Cook's first decade as CEO, 2021

Reuters: Asian supply chain after Huang, Nov 12

Korea JoongAng Daily: Korean chipmakers, Nov 12

IDC via BusinessWire: smartphones decline 0.5% in 2017

Epoch AI: frontier training compute grows 4-5x a year

Broadcom 8-K: Q3 FY26 results, Sep 2, 2026

Epoch AI: data on machine learning hardware

Jon Peddie Research: NVIDIA GTC 2026 keynote

NVIDIA proxy statement (DEF 14A), 2022: director biographies

NVIDIA blog: first DGX-1 delivered to OpenAI, Aug 15, 2016

CBS News: who will succeed Jensen Huang, Nov 12

Kalshi News: next NVIDIA CEO odds, Nov 12

Yahoo Finance: prediction markets eye next CEO, Nov 12

Bloomberg: OpenAI backs AI chip control measures, Nov 12

Semafor via Yahoo: Congress pushes tighter export controls

CNBC: Anthropic S-1 compute commitments, Oct 26

Cerebras 8-K: Q3 2026 results, Nov 12

CNBC: analysts weigh in on Nvidia's future, Nov 12

Motley Fool: Truist, Craig-Hallum downgrades, Nov 12

Bloomberg: Nvidia to buy Mellanox for $6.9B, Mar 11, 2019

Benzinga: early reaction to the Mellanox deal, Mar 2019

TSMC 6-K: Q3 2026 results, Oct 15

Intel 8-K: Q3 2026 results, Oct 29

CompaniesMarketCap: semiconductor P/E ratios, Nov 12

InvestingLive: chip stock forward P/Es, Nov 12

NVIDIA blog: Ian Buck author biography

Fortune: Ian Buck, most powerful rising executives 2025

Seoul Economic Daily: new TPUs, Helios, Nov 10

Press-reported: allegations, intraday moves, stake sizes, $530B tally, rival launches, HBM estimate · Photos supplied by the user