

AI SYSTEMS

Jev: A Model That Returns Decisions

By the end you'll be able to say what Jev returns, what is known about how it's built, and where it fits in an agent.

Explainer · September 2026

Where we're going

01 What Jev returns

Probabilities, not paragraphs

02 Why people care

Speed and cost, on TypeSafe's own numbers

03 How it's built

One named method, many unknowns

04 Where it fits

A fast decision layer in an agent

Jev returns a decision your code can act on, not a paragraph

- You send a state and a question with fixed answers.
- It returns a probability for each option.
- TypeSafe calls this a System One model.

Early access opened September 15, 2026. API only.

An LLM writes an answer; Jev picks one

LLM

Generates text token by token.

Your code must parse the reply.

Best for reasoning and writing.

Jev


Returns all outputs in one pass, per TypeSafe.

Answers arrive as probabilities.

Best for routing and approvals.

Match the output type to the job: text for writing, decisions for routing.

One pass turns a state and a question into probabilities



Your code compares the top bar to a threshold and acts.

Illustrative: bars show shape only, not real Jev output.

Jev answers three kinds of typed question



Choice

Pick one of the options

1 to 255 options



Score

Rate against ordered levels

2 to 10 levels



Noul

Yes or no

Probability, 0 to 1


Each type returns probabilities and a confidence.

Routing a support ticket takes four steps and no text parsing

- 1 A ticket arrives: "I was charged twice."
- 2 You define the options: billing, technical, sales, fraud.
- 3 Jev returns a probability for each option, plus a confidence.
- 4 Your code routes automatically above 0.90, otherwise asks a human.

The threshold is ordinary code, not a prompt.


In TypeSafe's own example, Jev took 0.114 s and the comparison LLM 8.566 s



About 75 times faster on this one query.

TypeSafe launch post. Vendor-reported; not independently reproduced.

In TypeSafe's example, Jev cost \$0.000081 per query; the comparison LLM \$0.013880



Jev's mean accuracy across TypeSafe's four workflows: 67.8%.

Vendor-reported (TypeSafe). Comparison model's accuracy not verified here.


WHY PEOPLE CARE

171×

Lower cost in the same example: \$0.000081 per query versus \$0.013880.

Vendor-reported. TypeSafe's broader claim is 40 to 400 times cheaper than frontier LLMs.


The person behind Jev worked on RLHF at OpenAI



- Worked on RLHF and InstructGPT at OpenAI.
- Now leads TypeSafe AI, the company behind Jev.
- RLHF is how chat models learn from human feedback.

Photo supplied by presenter. Background per TechCrunch and Wikipedia.

Reinforcement learning trains a model with rewards



The reward defines what the model learns to optimize.

HOW IT'S BUILT

RLCD

Reinforcement Learning for Calibrated Decisions: the training method TypeSafe says it used for Jev.

MODEL SAYS 0.80




RIGHT 8 OF 10 TIMES

Concept illustration, not Jev data.

WHY IT MATTERS

Calibrated means 0.80 should be right about 80% of the time, so thresholds in code can be trusted.

A calibration-aware training loop scores confidence, not just correctness



TypeSafe has not published RLCD's actual reward or algorithm.

A confident wrong answer costs far more than a hedged one

Truth: the ticket belongs to billing.

Confident and right

says billing 90%

0.02

Hedged and wrong

says fraud 60%

0.72

Confident and wrong

says fraud 90%

1.62

Penalty (Brier score, lower is better)


Hypothetical inputs, not TypeSafe's reward. Values computed from the Brier formula.

Most of how Jev is built has not been disclosed

Topic	What is public
Architecture	Claimed new, with a parallel sampler; details unpublished
Training data	Reported as synthetic only
Base model	Not disclosed; open-weight base suspected
Weights	Not published; API only


Reporting on the base model is unconfirmed.

Jev fits as a fast decision layer in front of heavier tools




Suggested pattern, not TypeSafe's published architecture.

Next: a live demo that triages support tickets with Jev



SAMPLE RUN: FIVE TICKETS



Output from the demo notebook's saved run.

What to take away

- 01 Jev returns probabilities over options you define, not free text.
- 02 Speed and cost gains are TypeSafe's own numbers, not yet independently reproduced.
- 03 RLCD aims for calibrated confidence; the base model and data are undisclosed.
- 04 Use it as a fast decision layer; keep LLMs for reasoning and writing.

Questions?

GO DEEPER

TypeSafe launch post: typesafe.ai/blog/introducing-system-one-models-and-jev

TechCrunch launch coverage, September 18, 2026

MarkTechPost technical write-up, September 19, 2026

Demo notebook: [jev_ticket_triage.ipynb](#) (Colab)

Explainer · September 2026