

AGENTIC AI @ UIUC

LECTURE

MODEL-CONTEXT-PROTOCOL (MCP)


Agenda

- | | |
|---|--------|
| 1. AI Market News | 5 min |
| 2. The Problem & The MCP Solution | 5 min |
| 3. Under the Hood → Architecture | 10 min |
| 4. How to Use It → Building an MCP Server | 10 min |
| 5. Limitations & Security | 5 min |
| 6. Future Improvements | 5 min |
| 7. Skills vs. MCP | 5 min |

Opus 4.7

Agentic coding performance by effort level


Anthropic internal autonomous agentic coding evaluation



OpenAI
@OpenAI

With computer use on macOS, Codex can now use any app by seeing, clicking, and typing with its own cursor.

It runs in the background without taking over your computer, working on tasks like frontend iteration, app testing, or any workflow that doesn't expose an API.



Forge Code

Search leaderboard		Select agents		Select models		Select organizations	
Rank	Agent	Model	Date	Agent Org	Model Org	Accuracy	
1	ForgeCode	GPT-5.4	2026-03-12	ForgeCode	OpenAI	81.8% ± 2.0	
2	ForgeCode	Claude Opus 4.6	2026-03-12	ForgeCode	Anthropic	81.8% ± 1.7	
3	TongAgents	Gemini 3.1 Pro	2026-03-13	BIGAI	Google	80.2% ± 2.6	
4	SageAgent	GPT-5.3-Codex	2026-03-13	OpenSage	OpenAI	78.4% ± 2.2	
5	ForgeCode	Gemini 3.1 Pro	2026-03-02	ForgeCode	Google	78.4% ± 1.8	
6	Droid	GPT-5.3-Codex	2026-02-24	Factory	OpenAI	77.3% ± 2.2	
7	Capy	Claude Opus 4.6	2026-03-12	Capy	Anthropic	75.3% ± 2.4	
8	Simple Codex	GPT-5.3-Codex	2026-02-06	OpenAI	OpenAI	75.1% ± 2.4	

PROBLEM



LLMs are Stateless


**APIs are traditional for modern
software development**

**APIs are structured for humans to understand
systems not agents**



Introducing MCP


Visual Representation of MCP



Under the Hood: Architecture

Client-Server Model

Application like Claude Desktop hosts the **Client**. You build the **Server**.



Roles of the Client and Server

MCP Host: The user-facing application (Cursor, Claude Desktop, custom agent). Responsible for UI, maintaining context, and deciding when tools are needed.

Client: The protocol engine embedded inside the Host.

- Manages the 1:1 connection lifecycle.
- Executes the initial capability negotiation.
- Formats and routes messages.

Server: A lightweight, reactive backend exposing specific capabilities.

- Sits adjacent to the data source (local, internet network, or cloud).
- Strictly stateless and reactive; it only responds to client requests.

The Three Core Primitives

Prompts (Instruction Architecture)

Reusable, parameterized templates. Used to standardize tasks like code reviews or email drafting.

Resources (Context Injection)

Read-only data payload exposed via URIs (file:// or some other endpoint like postgres://). Feeds content to the LLM without computational overhead.

Tools (Action Execution)

Arbitrary executable functions (@mcp.tool). Allows the LLM to interact with the outside world (API calls, db writes, file editing). Requires strict input schemas.

Custom Tool Integration



MCP essentially collapses the NxM matrix down to N+M (n clients/hosts, m external tools/data sources).

Build an MCP server once, and it instantly works on every compliant IDE and agent platform.

The Client Landscape

Chat Interfaces: Claude Desktop, ChatGPT

AI Native IDEs: Cursor, Windsurf, Zed, Replit

Autonomous Coding Agents: Cline / Roo Code

The Portability Benefit: Developers can establish a customized “toolchain” of MCP servers locally and easily hot-swap their frontend client or underlying model.




High-Utility MCP Servers

Databases: postgres and sqlite servers allow LLMs to inspect schemas and run read-only queries

File Systems & Git: Local filesystem servers with custom permissions and GitHub management.

Web Automation: playwright, puppeteer, mcp-chrome, browserbase

Enterprise Apps: slack, google-drive, etc.




Building your own MCP Server (Demo)

main.py 1 claude_desktop_config.json X

> Users > dansi > AppData > Local > Packages > Claude_pzs8srxjfjc > LocalCache > Roaming > Claude > claude_desktop_config.json > mcpServ

```
1  {
2    "mcpServers": {
3      "Demo" : {
4        "command": "uv",
5        "args": [
6          "run",
7          "--with",
8          "mcp[cli]",
9          "mcp",
10         "run",
11         "C:\\data\\dev\\agentikai\\PythonMCPServer\\mcp-server-demo\\main.py"
12       ]
13     }
14   }
15 }
```



D DEMO

Tool permissions


Choose when Claude is allowed to use these tools.

▼ Other tools 1

☑ Always allow ▼

add






Limitations & Security

The Context Bottleneck & Statelessness

MCP is stateless by design: MCP servers do not inherently remember past interactions. The host (client) must maintain the entire memory state.

Context window issues: Returning unstructured or massive payloads (like a huge CSV or JSON logs) will instantly exhaust the LLM's token limit.




The Security Nightmare

Prompt Injection: If an LLM is compromised via malicious input, it can misuse the tools you gave it access to.

The Sandbox Rule: Never give an MCP tool root access or unrestricted database write-access.

Human-in-the-Loop (HITL): Crucial for destructive actions or high-stakes environments.



MCPS VS SKILLS



RECAP ON SKILLS

SKILLS

Structured Instructions
and rules in simple
markdown files

They are static it is your
job to keep improving
them!!!!

They are basically
recipes

claude > skills > cleanup > cleanup.md > # Cleanup Skill

```
1  # Cleanup Skill
2  **Description:** Organize and clean up codebase structure before pushing to GitHub – arrange important code into proper folders, move auxiliary
   files, and ensure professional repository structure.
3  **Usage:** /cleanup
4
5  **Trigger this skill when:**
6  - User says "cleanup", "organize codebase", "prepare for GitHub"
7  - Before making a repository public
8  - Codebase has grown organically and needs structure
9  - Files are scattered in root directory
10 - Preparing for a major release or PR
11
12 **Skip for:** Already well-organized repos, single-file projects, documentation-only changes
13
14 ## Execution Workflow
15
16 ### Step 1: Analyze Current Structure
17 **Run these commands in parallel:**
18 - `find . -maxdepth 2 -type f -name "*.py" | head -30` – Python files
19 - `find . -maxdepth 2 -type f -name "*.js" -o -name "*.ts" | head -30` – JS/TS files
20 - `find . -maxdepth 1 -type f` – Root-level files
21 - `ls -la` – Full root directory listing
22 - `cat .gitignore 2>/dev/null || echo "No .gitignore"` – Check gitignore
23
24 **Identify:**
25 - Which files are in the root that shouldn't be (non-config code files)
26 - Existing folder structure patterns
```

Question 1

Do I need infra or just prompts?

MCP = heavy (servers, auth, infra)

Skills = lightweight (just text/patterns)

Question 2 - Context vs connectivity tradeoff?

MCP = solves integration problem

Skills = solves context overload problem



SO WTF TO DO???????

**WHEN SHOULD I USE MCP VS
SKILLS???**

**SHORT ANSWER : YOU NEED
BOTH!!!!**

When to use MCP?

Use MCPs when you need:

- Integrations
- Action
- Data is dynamic

MCPs help you connect with external providers, CRM, tools, etc.

When to use skills?

Use skills when you need

- You want fast iteration
- You're encoding workflows (e.g., presentation coaching logic)
- You don't need real-time external data

Skills are needed when you want to give the LLM structured inputs for repetitive tasks.