

BADM 372

From Venture Voice to Viva

By the end you'll be able to recognize when a client's feature request is really a system, and what it costs to generalize one.

Built by AI Hub · Information Systems Practicum

Where we're going

- | | | |
|-----------|------------------------|--------------------------------------|
| 01 | The problem | Customer discovery in an iMBA course |
| 02 | Venture Voice | Architecture, flow, and what it cost |
| 03 | What we learned | The reusable part wasn't the feature |
| 04 | Viva | More teacher control, many courses |
| 05 | Your practicum | What transfers to your client |

01

Customer discovery doesn't scale

Venture Voice began in an iMBA entrepreneurship course, where the assignment was easy to give and nearly impossible to supervise.

An iMBA entrepreneurship course had to teach interviewing at scale

- Customer discovery is the course's core practical skill.
- Real customers are scarce, slow, and available once.
- One instructor cannot observe an online cohort's interviews.

The bottleneck is feedback, not access to software.

THE REQUIREMENT

Three constraints, and dropping any one destroys the exercise

- Realistic — the customer must push back credibly.
- Repeatable — every student meets the same standard.
- At scale — no per-student instructor time.


Everything built afterward follows from refusing to drop one.

02

Building Venture Voice

Turning that requirement into personas, conversations, rubrics, and the data they leave behind.

One Next.js app handles auth, personas, scoring, and instructor reads



Real-time audio and video flow browser to vendor directly — the server brokers keys, never media.

Six layers, and only the middle one is ours to operate

Layer	Technology	Job
Client	Next.js / React	Live interview, cohort dashboard
Application	Next.js API on Azure	Auth, sessions, personas, scoring
Identity	Microsoft Entra ID	Campus SSO, instructor allowlist
AI services	Claude, ElevenLabs, Anam	Persona, voice, avatar
Data	Azure PostgreSQL + Prisma	User → Interview → Transcript + Score

Delivery: GitHub Actions — staging branch to slot, main to production.

One model plays two roles: the customer during, the grader after

Service	Acts as	Runs	What it does
Claude — persona	The customer	During the call	Stays in character, pushes back
Claude — scorer	The grader	After hang-up	Scores transcript to rubric
ElevenLabs	The voice	Audio interviews	Real-time speech both ways
Anam	The face	Video interviews	Lip-synced photoreal avatar

The rubric scores questioning, listening, and problem discovery.

Tracing one interview from sign-in to cohort dashboard

- 1 Student signs in with campus SSO and describes their startup idea.
- 2 The system builds an AI customer persona from that context.
- 3 Browser opens a live call — ElevenLabs voice or Anam avatar — Claude plays the customer.
- 4 On hang-up, the transcript is stored and scored against the rubric.

Results roll up into the student's dashboard and the instructor's cohort view.

Three choices keep the system small enough for one team to run

- One deployable unit — frontend, API, and orchestration ship together.
- Media flows browser to vendor; the server brokers keys only.
- Analytics computed at read time, so nothing aggregated can drift.

Interviews are the single source of truth — everything else is derived.

The AI was the easy part; the structure around it wasn't

- Personas had to be believable, not agreeable.
- Rubrics had to score the same way twice.
- Feedback had to name a specific next action.

Roughly true of every AI feature you will build.

One result validated the idea; the other forced the rewrite

What worked

Students repeated one interview five times.

Nobody does that with a real customer.

The value was practice, not grading.

What didn't

Each new assignment meant new code.

The system had no concept of an assignment.

Personas and rubrics were hard-coded.

The failure named the missing abstraction.

03

Venture Voice → Viva

Two gaps drove the rewrite: the teaching team had no way in, and the whole thing was welded to one course.

Venture Voice worked for one course and left its teachers outside

- Instructors could not shape personas, rubrics, or feedback.
- Mentors had no place in the review loop at all.
- A second course would have meant forking the codebase.

Viva was scoped to fix exactly these two problems.

Generalization

Rewriting a working feature so the parts that vary become configuration instead of code.

WHY IT MATTERS


Venture Voice served one course. Viva lets any instructor configure the assignment, any mentor review the results, and any course adopt it without new code.

Every hard-coded decision became something a course can configure

Was hard-coded	Became	Set by	Enables
The customer	Persona / examiner	Instructor	Any conversation partner
The question flow	Assignment type	Instructor	Interview, viva, defense
The grading logic	Rubric	Instructor	Course-specific standards
The saved output	Transcript + evidence	System	Review and analytics
Nobody reviewed	Instructor + mentor roles	Platform	Humans in the grading loop

One codebase, many courses — adopting it no longer requires engineering.

Configuration, conversation, and assessment form one closed loop



What instructors and mentors see at the end reshapes what they configure next term.

Same conversation engine, a different system around it

Venture Voice

One course, one assignment type.

Personas and rubrics fixed in code.

Teaching team saw results, shaped nothing.

Viva

Any course, many assignment types.

Personas and rubrics configured per course.

Instructors and mentors inside the review loop.

The interview didn't change. Who controls it did.

The same team built in two opposite directions

- Venture Voice: problem → product. A known gap, then a system.
- Viva: product → problem. A working system, then new gaps to fill.
- Both are legitimate. They fail in opposite ways.

Problem-first risks a tool nobody reuses; product-first risks a solution chasing a use.

Following one student interaction from setup to gradebook

- 1 Instructor picks an assignment type, a persona, and a rubric.
- 2 Student holds a live conversation inside Viva.
- 3 Viva stores the transcript and tags rubric evidence.
- 4 Scores generate; the instructor reviews, adjusts, and teaches from analytics.

One interaction, five artifacts, no instructor time until review.

What to take into your practicum

- 01 Your client asks for a feature. Find the system underneath it.
- 02 Start problem → product. Earn the right to go product → problem.
- 03 What you choose to record decides what you can ever learn.
- 04 Ship into real use; you cannot reason your way to what breaks.
- 05 Design for the people running the course, not only the students.

Questions?

GO DEEPER

dl-venture-voice.azurewebsites.net — the live system

Bring your client's first request to office hours